

The image shows the front cover of a spiral-bound notebook. The cover is a light beige or cream color with a subtle, repeating pattern of small, faint, stylized floral or geometric motifs. The spiral binding is visible on the left side, consisting of a series of metal loops. The text is centered on the cover.

CPT-S 415

Big Data

Yinghui Wu

EME B45

CPT-S 415

Big Data

Data dependencies and DBMS

- ✓ RDBMS design
 - the normal forms 1NF, 2NF, 3NF, BCNF
 - normal forms transformation
- ✓ DBMS: architecture

A spiral-bound notebook with a brown cover and a cream-colored page. The spiral binding is on the left side. A horizontal line is drawn across the page, and a grey rectangular box is positioned in the center.

Database Management Systems (DBMS)

Database management systems (DBMS)

- ✓ A **Database** is a collection of stored operational data used by the application systems of some particular enterprise (C.J. Date)
- ✓ A **DBMS** is a computer software with the capability to 1) store data in an integrated, structured format and to 2) enable users to retrieve, manipulate and manage the data.
 - Paper “Databases”
 - Still contain a large portion of the world’s knowledge
 - File-Based Data Processing Systems
 - Early batch processing of (primarily) business data
 - Database Management Systems (DBMS)

Data ≠ Information ≠ Knowledge!



Data



Information



Presentation



Knowledge



EpicGraphic.com

er is most
buy new

n is profitable
er and is likely
n carriers

Importance of DBMS

- It helps make data management more efficient and effective.
- Its query language allows quick answers to *ad hoc* queries.
- It provides end users better access to more and better-managed data.
- It promotes an integrated view of organization's operations -- “big picture.”
- It reduces the probability of inconsistent data.

RDBMS design




A Big data Fallacy

- ✓ “Database Design in the era of Big data is less important” (?)
 - New high-volume data streams
 - specialized hardware/software
 - Storage issues coped by hardware appliance
- ✓ **Fact:**
 - Most data is physically located in DBMS and new special-purpose appliance
 - Data loads, extract, transform, preprocessing ops continue as is
 - Database design for quality assurance

Big data a necessity at Largest Scale

A certain kind of developer at
a certain kind of company

Most development still RDBMS



MySQL, Oracle,
Mongo, Cassandra,
some memcache,
Some Hadoop...

Relational Databases Design

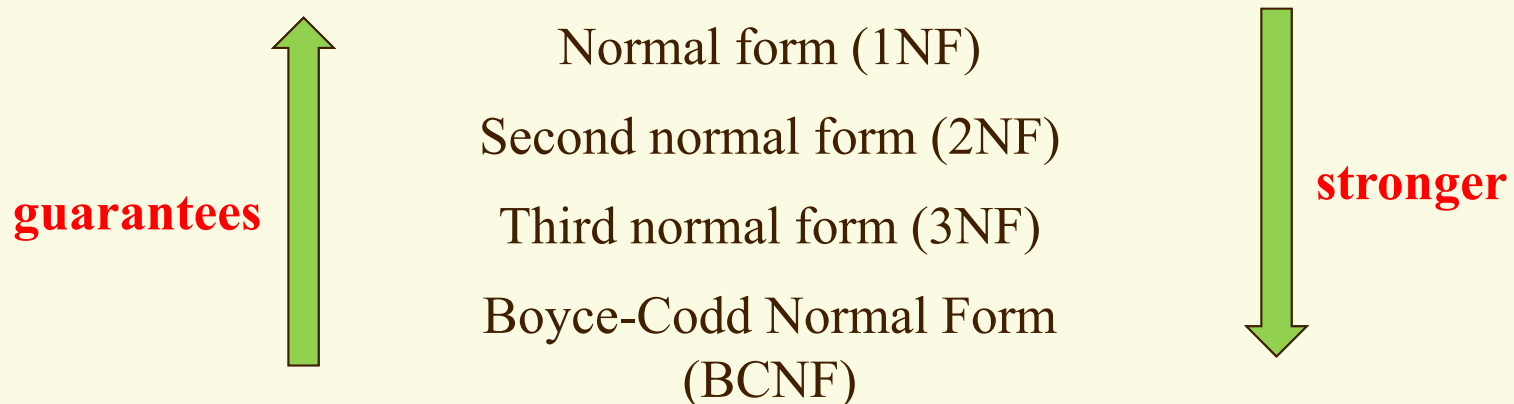
- ✓ Relational database design: The grouping of attributes to form "good" relation schemas
- ✓ Two levels of relation schemas:
 - The logical "user view" level
 - The storage "base relation" level
- ✓ Design is concerned mainly with base relations
- ✓ We have assumed schema R is given
 - R could have been generated when converting E-R diagram to a set of tables.
 - R could have been a single relation containing *all* attributes that are of interest (called **universal relation**).
 - **Normalization** breaks R into smaller relations.
 - R could have been the result of some ad hoc design of relations, which we then test/convert to normal form.

A spiral-bound notebook with a brown cover and a cream-colored page. A horizontal line is drawn across the page, and a gray rectangular box is positioned in the center. The text "Database Tables and Normalization" is written in red, italicized font within this box.

Database Tables and Normalization

Introduction to Normalization

- ✓ **Normalization:** Process of decomposing unsatisfactory "bad" relations by breaking up their attributes into smaller relations
- ✓ **Normal form:** Condition using keys and functional dependencies (FDs) of a relation to certify whether a relation schema is in a particular normal form
 - 2NF, 3NF, BCNF based on keys and FDs of a relation schema
 - 4NF based on keys, multi-valued dependencies



The Need for Normalization

- ✓ Mixing attributes of multiple entities may cause problems
 - Information is stored redundantly wasting storage
 - Problems with update anomalies:
 - Insertion anomalies
 - Deletion anomalies
 - Modification anomalies
- ✓ The report may yield different results, depending on data anomaly
 - Primary keys?
 - Data redundancy
 - Possible data inconsistencies: JOB_CLASS: Elect.Engineer, Elect.Eng. EI.Eng. EE

The Need for Normalization

- ✓ Example: company that manages building projects
 - Charges its clients by billing hours spent on each contract
 - Hourly billing rate is dependent on employee's position

TABLE 5.1 A SAMPLE REPORT LAYOUT

PROJ. NUM.	PROJECT NAME	EMPLOYEE NUMBER	EMPLOYEE NAME	JOB CLASS.	CHG/HOUR	HOURS BILLED	TOTAL CHARGE
15	Evergreen	103	June E. Arbough	Elec. Engineer	\$84.50	23.8	\$2,011.10
		101	John G. News	Database Designer	\$105.00	19.4	\$2,037.00
		105	Alice K. Johnson *	Database Designer	\$105.00	35.7	\$3,748.50
		106	William Smithfield	Programmer	\$35.75	12.6	\$450.45
		102	David H. Senior	Systems Analyst	\$96.75	23.8	\$2,302.65
Subtotal							\$10,549.70
18	Amber Wave	114	Annelise Jones	Applications Designer	\$48.10	24.6	\$1,183.26
		118	James J. Frommer	General Support	\$18.36	45.3	\$831.71
		104	Anne K. Ramoras *	Systems Analyst	\$96.75	32.4	\$3,135.70
		112	Darlene M. Smithson	DSS Analyst	\$45.95	44.0	\$2,021.80
Subtotal							\$7,171.47
22	Rolling Tide	105	Alice K. Johnson	Database Designer	\$105.00	64.7	\$6,793.50
		104	Anne K. Ramoras	Systems Analyst	\$96.75	48.4	\$4,682.70
		113	Delbert K. Joenbrood *	Applications Designer	\$48.10	23.6	\$1,135.16
		111	Geoff B. Wabash	Clerical Support	\$26.87	22.0	\$591.14
		106	William Smithfield	Programmer	\$35.75	12.8	\$457.60
Subtotal							\$13,660.10
25	Starflight	107	Maria D. Alonzo	Programmer	\$35.75	24.6	\$879.45
		115	Travis B. Bawangi	Systems Analyst	\$96.75	45.8	\$4,431.15
		101	John G. News *	Database Designer	\$105.00	56.3	\$5,911.50
		114	Annelise Jones	Applications Designer	\$48.10	33.1	\$1,592.11
		108	Ralph B. Washington	Systems Analyst	\$96.75	23.6	\$2,283.30
		118	James J. Frommer	General Support	\$18.36	30.5	\$559.98
		112	Darlene M. Smithson	DSS Analyst	\$45.95	41.4	\$1,902.33
Subtotal							\$17,559.82
Total							\$48,941.09
* Indicates project leader							

* Indicates project leader

A Table in the Report Format

Table name: RPT_FORMAT

Database name: Ch05_ConstructCo

	PROJ_NUM	PROJ_NAME	EMP_NUM	EMP_NAME	JOB_CLASS	CHG_HOUR	HOURS
▶	15	Evergreen	103	June E. Arbough	Elect. Engineer	\$84.50	23.8
			101	John G. News	Database Designer	\$105.00	19.4
			105	Alice K. Johnson *	Database Designer	\$105.00	35.7
			106	William Smithfield	Programmer	\$35.75	12.6
			102	David H. Senior	Systems Analyst	\$96.75	23.8
	18	Amber Wave	114	Annelise Jones	Applications Designer	\$48.10	24.6
			118	James J. Frommer	General Support	\$18.36	45.3
			104	Anne K. Ramoras *	Systems Analyst	\$96.75	32.4
			112	Darlene M. Smithson	DSS Analyst	\$45.95	44.0
	22	Rolling Tide	105	Alice K. Johnson	Database Designer	\$105.00	64.7
			104	Anne K. Ramoras	Systems Analyst	\$96.75	48.4
			113	Delbert K. Joenbrood *	Applications Designer	\$48.10	23.6
			111	Geoff B. Wabash	Clerical Support	\$26.87	22.0
			106	William Smithfield	Programmer	\$35.75	12.8
	25	Starflight	107	Maria D. Aristo	Programmer	\$35.75	24.6
			115	Travis B. Bawangi	Systems Analyst	\$96.75	45.8
			101	John G. News *	Database Designer	\$105.00	56.3
			114	Annelise Jones	Applications Designer	\$48.10	33.1
			108	Ralph B. Washington	Systems Analyst	\$96.75	23.6
			118	James J. Frommer	General Support	\$18.36	30.5
			112	Darlene M. Smithson	DSS Analyst	\$45.95	41.4

The Need for Normalization

- ✓ For most business database design, 3NF is highest we need to go in the normalization process
 - Highest level of normalization is not always most desirable
- ✓ Normalization should be part of design process
 - Make sure that proposed entities meet required normal form before table structures are created
 - Many real-world databases have been improperly designed or burdened with anomalies if improperly modified during course of time
 - You may be asked to redesign and modify existing databases

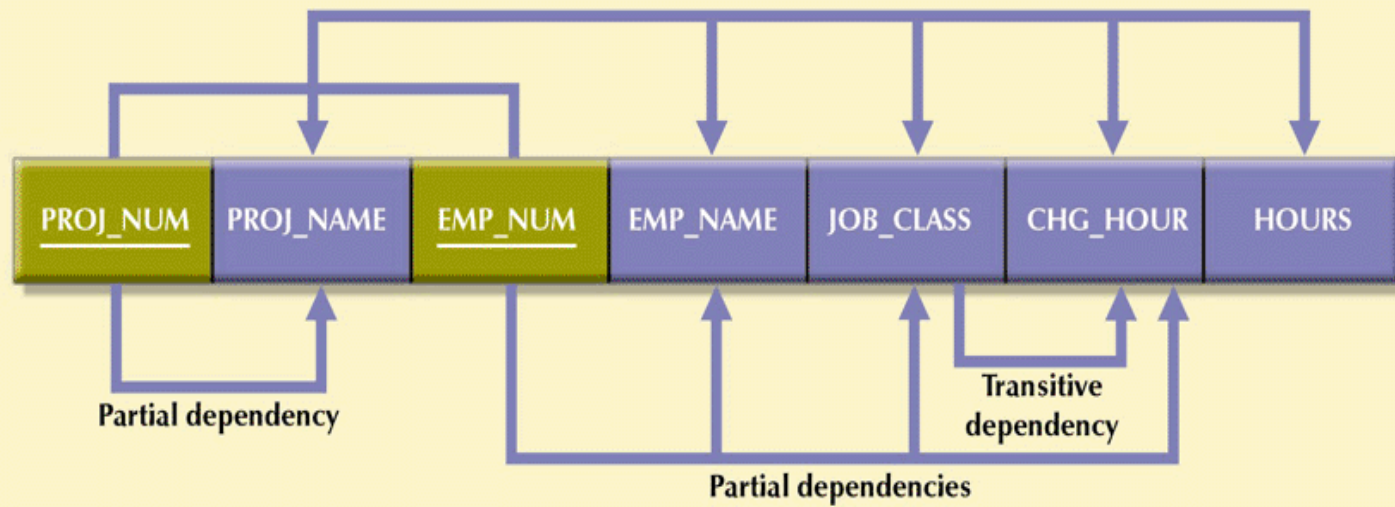
Functional Dependencies

- ✓ **Functional dependencies** (FDs) are used to specify formal measures of the "goodness" of relational designs
- ✓ FDs and keys are used to define normal forms for relations
- ✓ FDs are constraints that are derived from the meaning and interrelationships of the data attributes

Functional Dependencies (2)

- ✓ A set of attributes X *functionally determines* a set of attributes Y if the value of X determines a unique value for Y
- ✓ $X \rightarrow Y$ holds if whenever two tuples have the same value for X , they *must have* the same value for Y
If $t_1[X]=t_2[X]$, then $t_1[Y]=t_2[Y]$ in any relation instance $r(R)$
- ✓ a *constraint* on **all** relation instances $r(R)$
- ✓ FDs are derived from the real-world constraints on the attributes
 - $SSN \rightarrow ENAME$
 - $PNUMBER \rightarrow \{PNAME, PLOCATION\}$
 - $\{SSN, PNUMBER\} \rightarrow HOURS$
 - what if LHS attributes is a **key**?
- ✓ A FD $X \rightarrow Y$ is a
 - **Full functional dependency** if removal of any attribute from X means the FD does not hold any more
 - **Transitive functional dependency**- if there a set of attributes Z that are neither a primary or candidate key and both $X \rightarrow Z$ and $Z \rightarrow Y$ holds.

A Dependency Diagram



Inference Rules for FDs

- ✓ Given a set of FDs F , we can infer additional FDs that hold whenever the FDs in F hold
- ✓ Armstrong's inference rules
 - A1. (*Reflexive*) If Y subset-of X , then $X \rightarrow Y$
 - A2. (*Augmentation*) If $X \rightarrow Y$, then $XZ \rightarrow YZ$
(Notation: XZ stands for $X \cup Z$)
 - A3. (*Transitive*) If $X \rightarrow Y$ and $Y \rightarrow Z$, then $X \rightarrow Z$
- ✓ A1, A2, A3 form a *sound and complete* set of inference rules

First Normal Form

- ✓ Tabular format in which:
 - All key attributes are defined
 - There are no repeating groups in the table
 - All attributes are dependent on primary key
- ✓ Disallows composite attributes, multivalued attributes, and **nested relations**; attributes whose values *for an individual tuple* are non-atomic

1NF deals with the “shape” of the tables

Transform to 1NF

- ✓ Step 1: Eliminate repeating groups
- ✓ Step 2: Identify primary key
- ✓ (Step 3: Identify all dependencies)

FIGURE 5.2 DATA ORGANIZATION: FIRST NORMAL FORM

Table name: DATA_ORG_1NF				Database name: Ch05_ConstructCo			
	PROJ_NUM	PROJ_NAME	EMP_NUM	EMP_NAME	JOB_CLASS	CHG_HOUR	HOURS
►	15	Evergreen	103	June E. Arbough	Elect. Engineer	\$84.50	23.8
	15	Evergreen	101	John G. News	Database Designer	\$105.00	19.4
	15	Evergreen	105	Alice K. Johnson *	Database Designer	\$105.00	35.7
	15	Evergreen	106	William Smithfield	Programmer	\$35.75	12.6
	15	Evergreen	102	David H. Senior	Systems Analyst	\$96.75	23.8
	18	Amber Wave	114	Annelise Jones	Applications Designer	\$48.10	24.6
	18	Amber Wave	118	James J. Frommer	General Support	\$18.36	45.3
	18	Amber Wave	104	Anne K. Ramoras *	Systems Analyst	\$96.75	32.4
	18	Amber Wave	112	Darlene M. Smithson	DSS Analyst	\$45.95	44.0
	22	Rolling Tide	105	Alice K. Johnson	Database Designer	\$105.00	64.7
	22	Rolling Tide	104	Anne K. Ramoras	Systems Analyst	\$96.75	48.4
	22	Rolling Tide	113	Delbert K. Joenbrood *	Applications Designer	\$48.10	23.6
	22	Rolling Tide	111	Geoff B. Wabash	Clerical Support	\$26.87	22.0
	22	Rolling Tide	106	William Smithfield	Programmer	\$35.75	12.8
	25	Starflight	107	Maria D. Alonzo	Programmer	\$35.75	24.6
	25	Starflight	115	Travis B. Bawangi	Systems Analyst	\$96.75	45.8
	25	Starflight	101	John G. News *	Database Designer	\$105.00	56.3
	25	Starflight	114	Annelise Jones	Applications Designer	\$48.10	33.1
	25	Starflight	108	Ralph B. Washington	Systems Analyst	\$96.75	23.6
	25	Starflight	118	James J. Frommer	General Support	\$18.36	30.5
	25	Starflight	112	Darlene M. Smithson	DSS Analyst	\$45.95	41.4

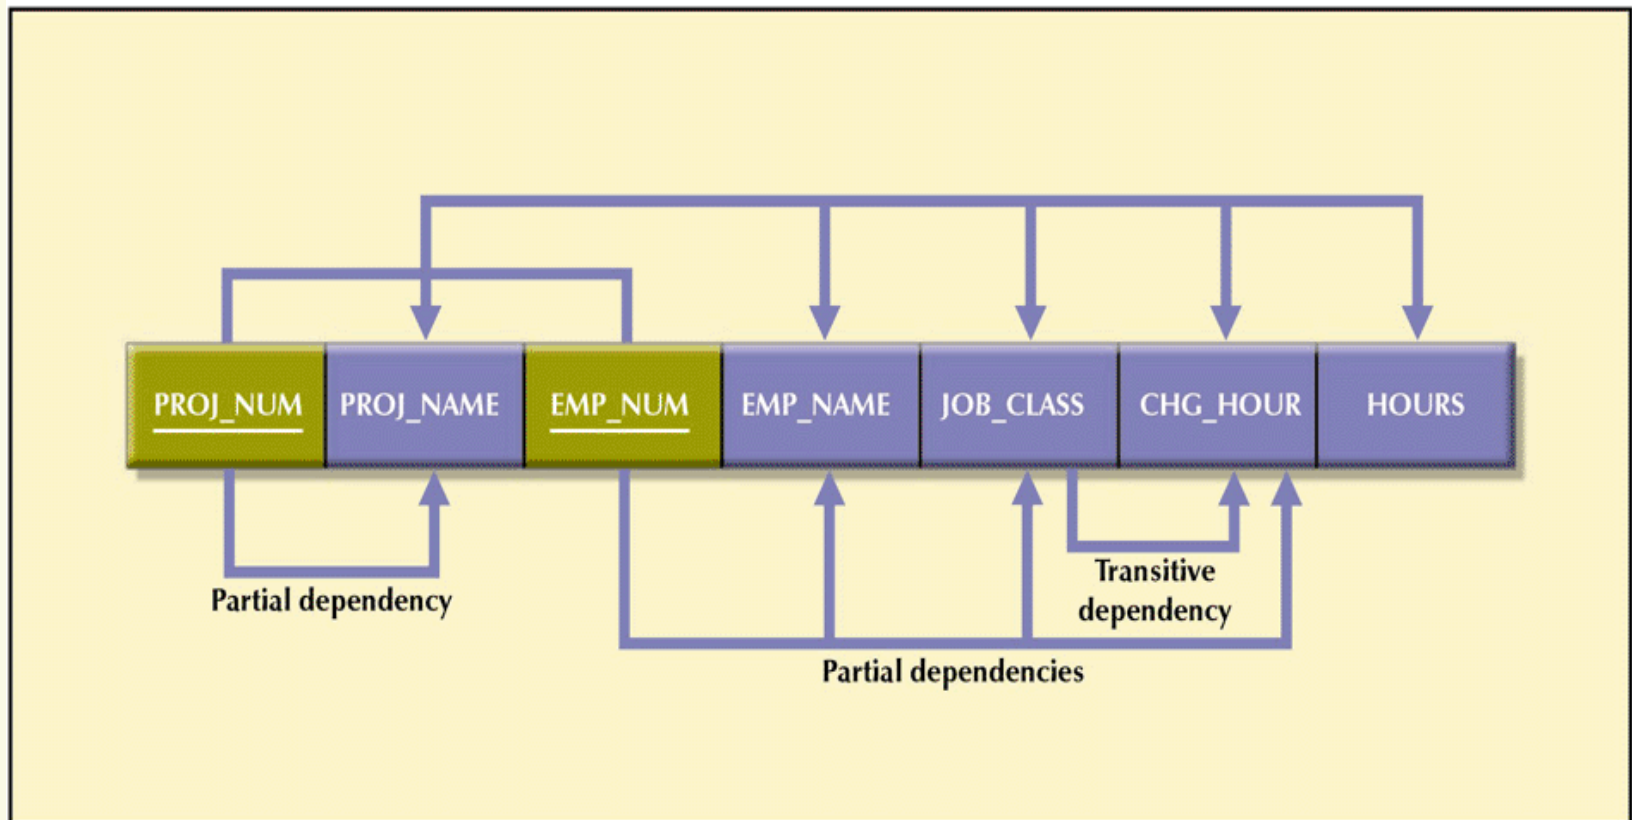
Second Normal Form

- ✓ Table is in second normal form (2NF) if:
 - It is in 1NF and
 - It includes no partial dependencies:
 - No attribute is dependent on only a portion of the primary key = every attribute A not in PK is fully functionally dependent on PK

2NF deals with the relationship between non-key and key attributes

A Dependency Diagram: First Normal Form (1NF)

FIGURE 5.3 A DEPENDENCY DIAGRAM: FIRST NORMAL FORM (1NF)

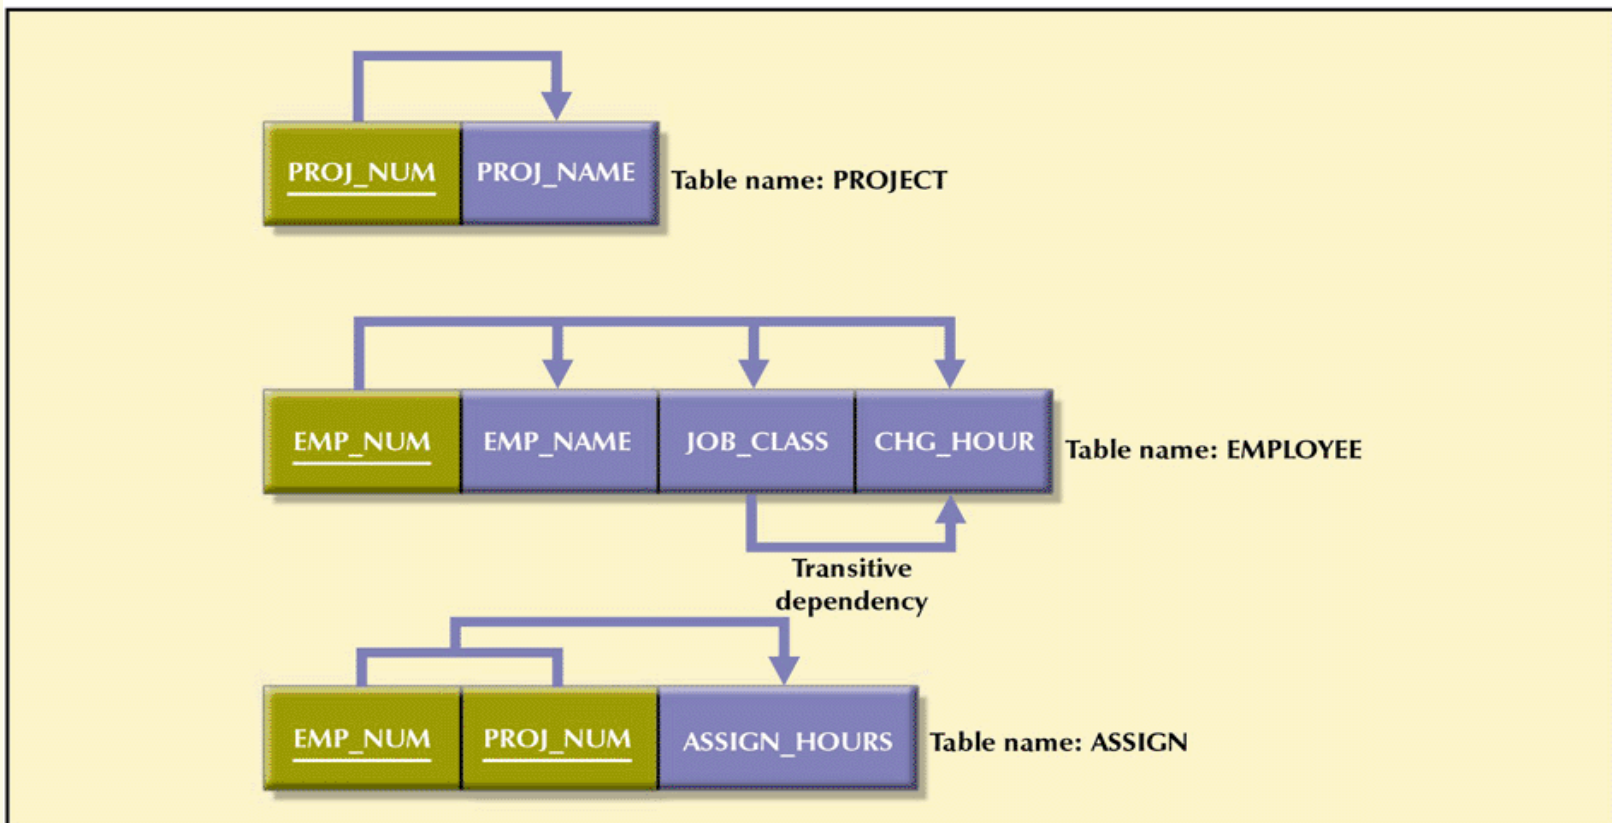


Conversion to Second Normal Form

- ✓ Relational database design can be improved by converting the database into second normal form (2NF)
- ✓ Two steps
 - Step 1: Write each key attribute on separate line, and then write the original (composite) key on the last line; Each component will become the key in a new table
 - Step 2: Determine which attributes are dependent on which other attributes (remove anomalies)

Second Normal Form (2NF) Conversion Results

FIGURE 5.4 SECOND NORMAL FORM (2NF) CONVERSION RESULTS



Third Normal Form

- ✓ A table is in third normal form (3NF) if:
 - It is in 2NF and
 - It contains no transitive dependencies

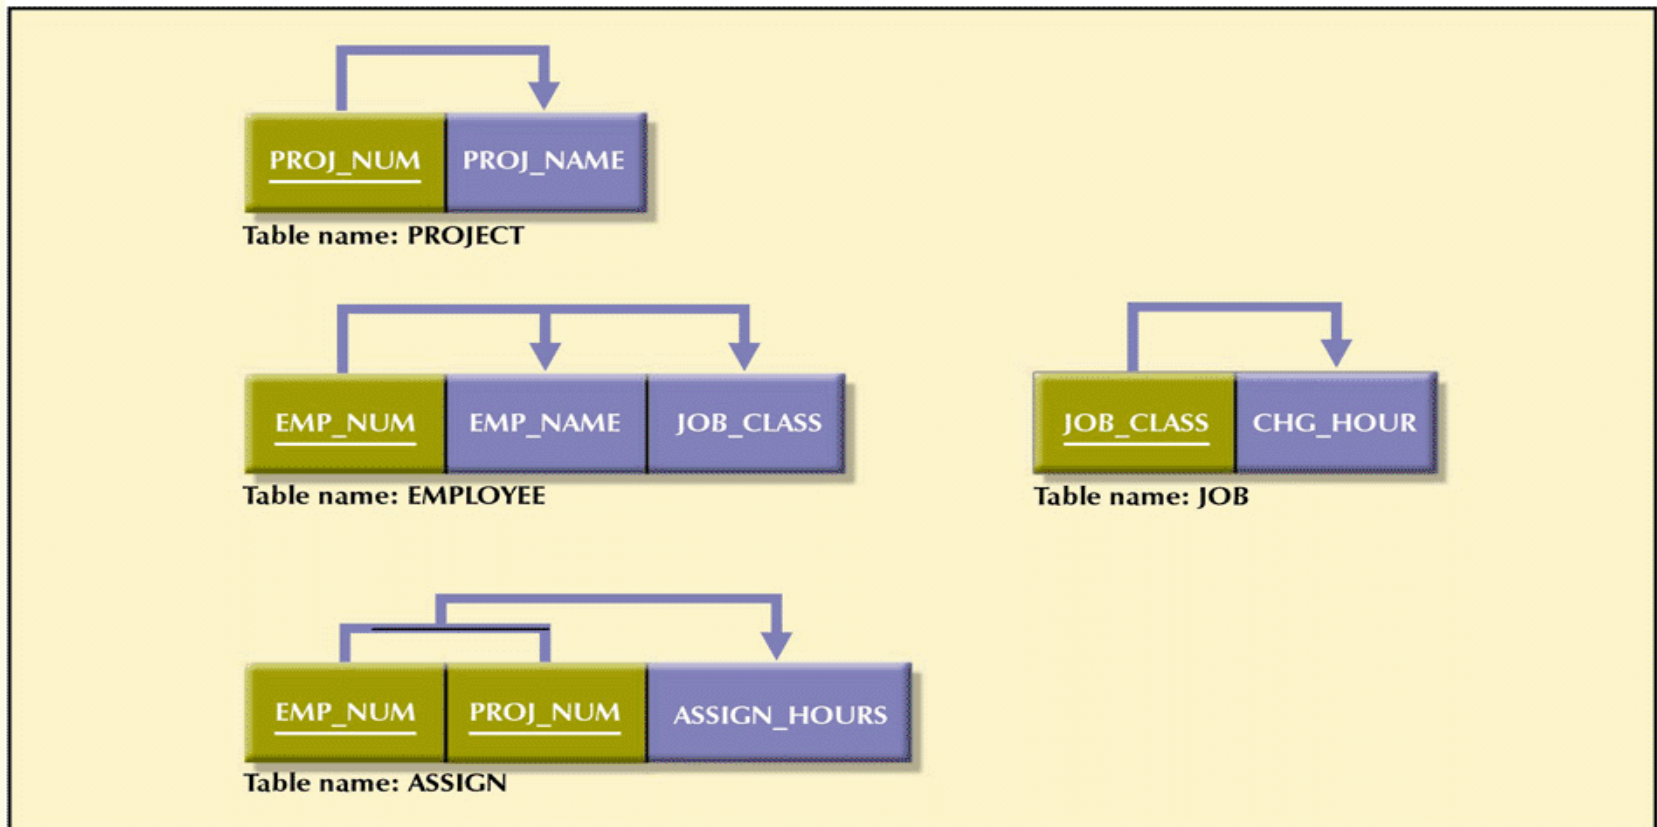
3NF removes transitive dependencies

Conversion to Third Normal Form

- ✓ Data anomalies created are easily eliminated by completing three steps
 - **Step 1: find new fact:** For every transitive dependency $X \rightarrow Y$, write fact Z as a PK for a new table where $X \rightarrow Z$ and $Z \rightarrow Y$
 - **Step 2: Identify the Dependent Attributes** - Identify the attributes dependent on each Z identified in Step 1 and identify the dependency; Name the table to reflect its contents and function
 - **Step 3: Remove $X \rightarrow Y$ from Transitive Dependencies**

Third Normal Form (3NF) Conversion Results

FIGURE 5.5 THIRD NORMAL FORM (3NF) CONVERSION RESULTS




*data depends on the key [1NF], the whole key [2NF]
and nothing but the key [3NF]*

BCNF (Boyce-Codd Normal Form)

- ✓ A relation schema R is in **Boyce-Codd Normal Form (BCNF)**, aka 3.5NF, if whenever an FD $X \rightarrow A$ holds in R, then X is a superkey of R
 - Each normal form is strictly stronger than the previous one:
 - Every 2NF relation is in 1NF
 - Every 3NF relation is in 2NF
 - Every BCNF relation is in 3NF
 - There exist relations that are in 3NF but not in BCNF
 - The goal is to have each relation in BCNF (or 3NF)
 - Why not 4NF? 4 and 5NF deal with multi-valued attributes
 - Just slightly stricter than 3NF

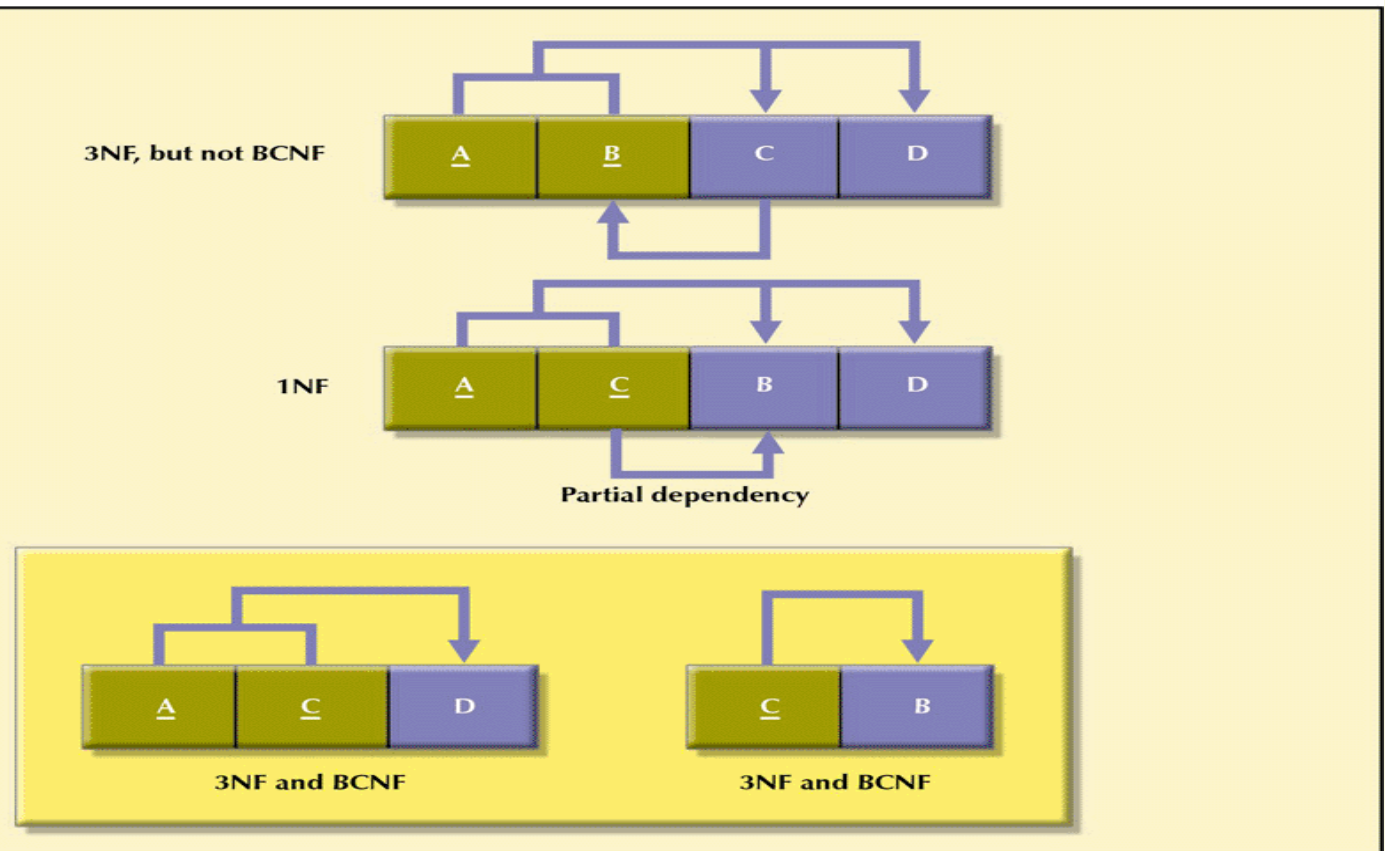
A Table That is in 3NF but not in BCNF

FIGURE 5.7 A TABLE THAT IS IN 3NF BUT NOT IN BCNF





Decomposition to BCNF

FIGURE 5.8 DECOMPOSITION TO BCNF



The Completed Database

FIGURE 5.6 THE COMPLETED DATABASE



Database name: Ch05_ConstructCo

Table name: PROJECT

	PROJ_NUM	PROJ_NAME	EMP_NUM
▶ +	15	Evergreen	105
+	18	Amber Wave	104
+	22	Rolling Tide	113
+	25	Starflight	101

Table name: JOB

	JOB_CODE	JOB_DESCRIPTION	JOB_CHG_HOUR
▶ +	500	Programmer	\$35.75
+	501	Systems Analyst	\$96.75
+	502	Database Designer	\$105.00
+	503	Electrical Engineer	\$84.50
+	504	Mechanical Engineer	\$67.90
+	505	Civil Engineer	\$55.78
+	506	Clerical Support	\$26.87
+	507	DSS Analyst	\$45.95
+	508	Applications Designer	\$48.10
+	509	Bio Technician	\$34.55
+	510	General Support	\$18.36

FIGURE 5.6 THE COMPLETED DATABASE (CONTINUED)

Table name: ASSIGN

ASSIGN_NUM ASSIGN_DATE PROJ_NUM EMP_NUM ASSIGN_HOURS ASSIGN_CHG_HOUR ASSIGN_CHARGE

Table name: ASSIGN

Database name: Ch05_ConstructCo

	ASSIGN_NUM	ASSIGN_DATE	PROJ_NUM	EMP_NUM	ASSIGN_HOURS	ASSIGN_CHG_HOUR	ASSIGN_CHARGE
▶	1001	04-Mar-04	15	103	2.6	\$84.50	\$219.70
	1002	04-Mar-04	18	118	1.4	\$18.36	\$25.70
	1003	05-Mar-04	15	101	3.6	\$105.00	\$378.00
	1004	05-Mar-04	22	113	2.5	\$48.10	\$120.25
	1005	05-Mar-04	15	103	1.9	\$84.50	\$160.55
	1006	05-Mar-04	25	115	4.2	\$96.75	\$406.35
	1007	05-Mar-04	22	105	5.2	\$105.00	\$546.00
	1008	05-Mar-04	25	101	1.7	\$105.00	\$178.50
	1009	05-Mar-04	15	105	2.0	\$105.00	\$210.00
	1010	06-Mar-04	15	102	3.8	\$96.75	\$367.65
	1011	06-Mar-04	22	104	2.6	\$96.75	\$251.55
	1012	06-Mar-04	15	101	2.3	\$105.00	\$241.50
	1013	06-Mar-04	25	114	1.8	\$48.10	\$86.58
	1014	06-Mar-04	22	111	4.0	\$26.87	\$107.48
	1015	06-Mar-04	25	114	3.4	\$48.10	\$163.54
	1016	06-Mar-04	18	112	1.2	\$45.95	\$55.14
	1017	06-Mar-04	18	118	2.0	\$18.36	\$36.72
	1018	06-Mar-04	18	104	2.6	\$96.75	\$251.55
	1019	06-Mar-04	15	103	3.0	\$84.50	\$253.50
	1020	07-Mar-04	22	105	2.7	\$105.00	\$283.50
	1021	08-Mar-04	25	108	4.2	\$96.75	\$406.35
	1022	07-Mar-04	25	114	5.8	\$48.10	\$278.98
	1023	07-Mar-04	22	106	2.4	\$35.75	\$85.80

Table name: EMPLOYEE

EMP_NUM EMP_LNAME EMP_FNAME EMP_INITIAL EMP_HIREDATE JOB_CODE

Table name: EMPLOYEE

	EMP_NUM	EMP_LNAME	EMP_FNAME	EMP_INITIAL	EMP_HIREDATE	JOB_CODE
▶	101	News	John	G	08-Nov-98	502
+	102	Senior	David	H	12-Jul-87	501
+	103	Arbough	June	E	01-Dec-94	503
+	104	Ramoras	Anne	K	15-Nov-85	501
+	105	Johnson	Alice	K	01-Feb-91	502
+	106	Smithfield	vWilliam		22-Jun-03	500
+	107	Alonzo	Maria	D	10-Oct-91	500
+	108	vWashington	Ralph	B	22-Aug-89	501
+	109	Smith	Larry	vV	18-Jul-95	501
+	110	Olenko	Gerald	A	11-Dec-93	505
+	111	vVabash	Geoff	B	04-Apr-89	506
+	112	Smithson	Darlene	M	23-Oct-92	507
+	113	Joebrood	Delbert	K	15-Nov-94	508
+	114	Jones	Annelise		20-Aug-91	508
+	115	Bawangi	Travis	B	25-Jan-90	501
+	116	Pratt	Gerald	L	05-Mar-95	510
+	117	vWilliamson	Angie	H	19-Jun-94	509
+	118	Frommer	James	J	04-Jan-04	510

The background of the slide is a spiral-bound notebook with a brown cover and a cream-colored page. A silver spiral binding is visible on the left side. A horizontal line is drawn across the page, and a grey rectangular box is positioned in the center.

DBMS architecture: overview

DBMS architectures

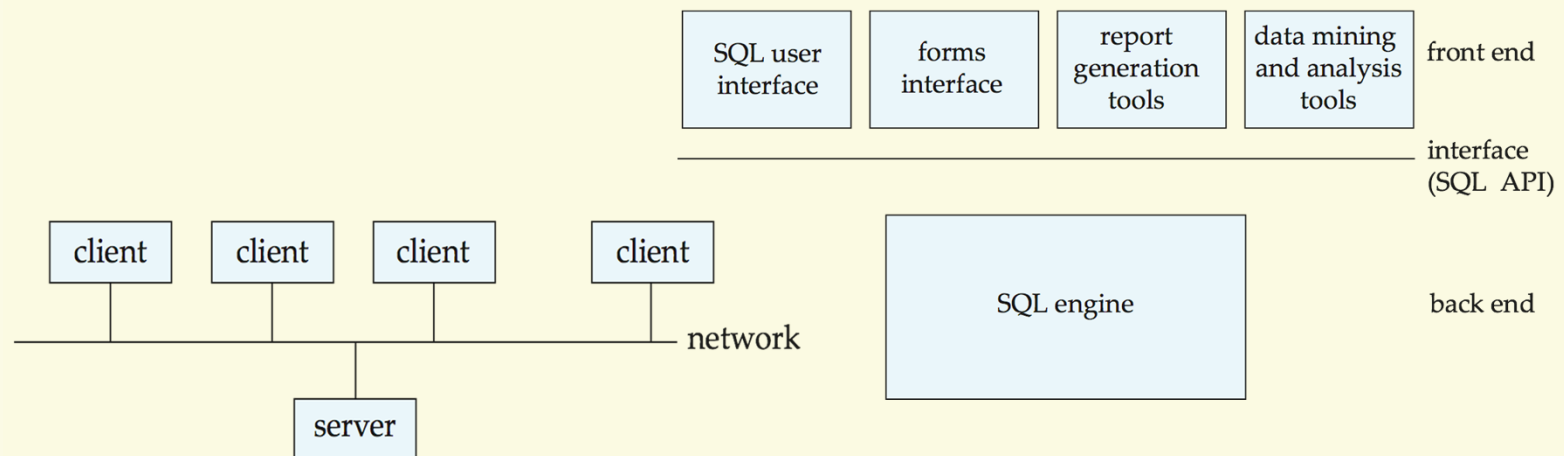
- ✓ Centralized and Client-Server Systems
- ✓ Parallel Systems
- ✓ Distributed Systems

Centralized Systems

- ✓ Run on a single computer system and do not interact with other computer systems.
 - Single-user system (e.g., personal computer or workstation): desk-top unit, single user, usually has only one CPU and one or two hard disks; the OS may support only one user.
 - Multi-user system: more disks, more memory, multiple CPUs, and a multi-user OS. Serve a large number of users who are connected to the system via terminals. Often called *server* systems.

Client-Server Systems

- ✓ Database functionality can be divided into:
 - **Back-end**: manages access structures, query evaluation and optimization, concurrency control and recovery.
 - **Front-end**: consists of tools such as *forms*, *report-writers*, and graphical user interface facilities.
- ✓ The interface between the front-end and the back-end is through SQL or through an application program interface.



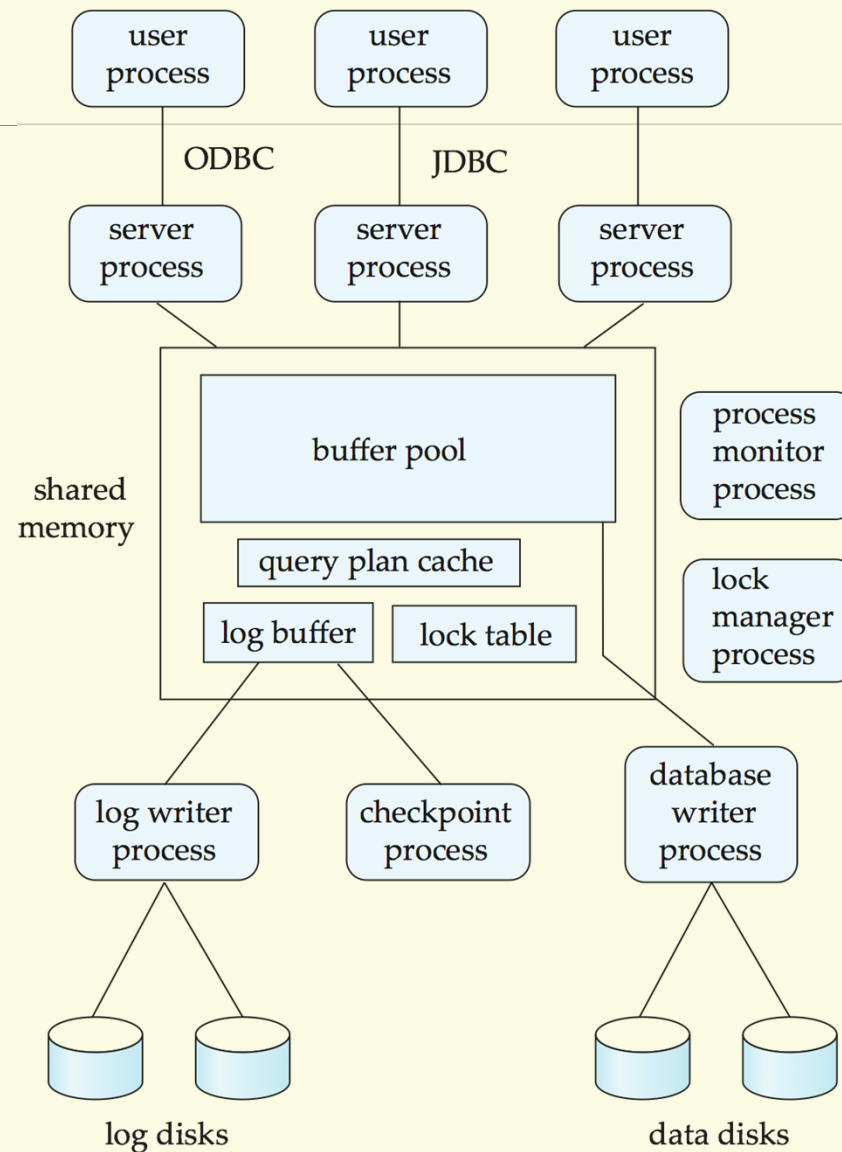
Server System Architecture

- ✓ Server systems can be broadly categorized into :
 - **transaction servers** which are widely used in relational database systems, and
 - **data servers**, used in object-oriented database systems

Transaction Servers

- ✓ Also called **query server** systems
 - Clients send requests to the server
 - Transactions are executed at the server
 - Results are shipped back to the client.
- ✓ Requests are specified in e.g., SQL, and communicated to the server through a *remote call* mechanism..
- ✓ *Open Database Connectivity* (ODBC) is a C language application program interface standard from Microsoft for connecting to a server, sending SQL requests, and receiving results.
- ✓ JDBC standard is similar to ODBC, for Java

Transaction System Processes (Cont.)



Data Servers

- ✓ Used in high-speed LANs, in cases where
 - The clients are comparable in processing power to the server
 - The tasks to be executed are compute intensive.
- ✓ Data are shipped to clients where processing is performed, and then shipped results back to the server.
- ✓ This architecture requires full back-end functionality at the clients.
- ✓ Used in many object-oriented database systems


Parallel Systems

- ✓ Parallel database systems consist of multiple processors and multiple disks connected by a fast interconnection network.
- ✓ A **coarse-grain parallel** machine consists of a small number of powerful processors
- ✓ A **massively parallel** or **fine grain parallel** machine utilizes thousands of smaller processors.
- ✓ Two main performance measures:
 - **throughput** --- the number of tasks that can be completed in a given time interval
 - **response time** --- the amount of time it takes to complete a single task from the time it is submitted


Speed-Up and Scale-Up

- ✓ **Speedup**: a fixed-sized problem executing on a small system is given to a system which is N -times larger.
 - Measured by:
$$\text{speedup} = \frac{\text{small system elapsed time}}{\text{large system elapsed time}}$$
 - Speedup is **linear** if equation equals N .
- ✓ **Scaleup**: increase the size of both the problem and the system
 - N -times larger system used to perform N -times larger job
 - Measured by:
$$\text{scaleup} = \frac{\text{small system small problem elapsed time}}{\text{big system big problem elapsed time}}$$
 - Scale up is **linear** if equation equals 1.


Speedup



Scaleup




Interconnection Architectures




(a) bus

send data on and receive data from a single communication bus; Does not scale well with increasing parallelism.



(b) mesh


each connected to all adjacent components; scales better. But may require $2\sqrt{n}$ hops to send message to a node




(c) hypercube

numbered in binary; connected to one another if binary differ in exactly 1 bit; n components connected to $\log(n)$ other components. can reach each other via at most $\log(n)$ links


Parallel Database Architectures




(a) shared memory



(b) shared disk




(c) shared nothing



(d) hierarchical


Shared Memory

- ✓ Processors and disks have access to a common memory, typically via a bus or through an interconnection network.
- ✓ Extremely efficient communication between processors — data in shared memory can be accessed by any processor without having to move it using software.
- ✓ Downside – architecture is not scalable beyond 32 or 64 processors since the bus or the interconnection network becomes a bottleneck
- ✓ Widely used for lower degrees of parallelism (4 to 8).



Shared Disk


- ✓ All processors can directly access all disks via an interconnection network, but the processors have private memories.
 - The memory bus is not a bottleneck
 - Architecture provides a degree of **fault-tolerance**
- ✓ Downside: bottleneck now occurs at interconnection to the disk subsystem.
- ✓ Shared-disk systems can scale to a somewhat larger number of processors, but communication between processors is slower.



IBM Sysplex and DEC clusters (now part of Compaq) running Rdb (now Oracle Rdb)

Shared Nothing

- ✓ Node consists of a processor, memory, and one or more disks. Processors at one node communicate with another processor at another node using an interconnection network.
- ✓ minimizing the interference of resource sharing.
- ✓ Shared-nothing multiprocessors can be scaled up to thousands of processors without interference.
- ✓ Main drawback: cost of communication and non-local disk access; sending data involves software interaction at both ends.




Teradata, Tandem,
Oracle-n CUBE

Hierarchical

- ✓ Combines characteristics of shared-memory, shared-disk, and shared-nothing architectures.
- ✓ Top level is a shared-nothing architecture – nodes connected by an interconnection network, and do not share disks or memory with each other.
- ✓ Each node of the system could be a shared-memory system with a few processors.
- ✓ Alternatively, each node could be a shared-disk system, and each of the systems sharing a set of disks could be a shared-memory system.
- ✓ Reduce the complexity of programming such systems by **distributed virtual-memory** architectures
 - Also called **non-uniform memory architecture (NUMA)**

Distributed Systems

- ✓ Data spread over multiple machines (also referred to as **sites** or **nodes**).
- ✓ Network interconnects the machines
- ✓ Data shared by users on multiple machines



Distributed Databases

- ✓ Homogeneous distributed databases
 - Same software/schema on all sites, data may be partitioned among sites
 - Goal: provide a view of a single database, hiding details of distribution
- ✓ Heterogeneous distributed databases
 - Different software/schema on different sites
 - Goal: integrate existing databases to provide useful functionality
- ✓ Differentiate between *local* and *global* transactions
 - A **local transaction** accesses data in the *single* site at which the transaction was initiated.
 - A **global transaction** either accesses data in a site different from the one at which the transaction was initiated or accesses data in several different sites.

Trade-offs in Distributed Systems

- ✓ Sharing data – users at one site able to access the data residing at some other sites.
- ✓ Autonomy – each site is able to retain a degree of control over data stored locally.
- ✓ Higher system availability through redundancy — data can be replicated at remote sites, and system can function even if a site fails.
- ✓ Disadvantage: added complexity required to ensure proper coordination among sites.
 - Software development cost.
 - Greater potential for bugs.
 - Increased processing overhead.

Summary

- ✓ RDBMS design process
 - Normalization
 - the normal forms 1NF, 2NF, 3NF, BCNF, and 4NF
 - normal forms transformation
- ✓ DBMS architectures

