CPT-S 415

Big Data

Yinghui Wu EME B45

CPT-S 415 Big Data

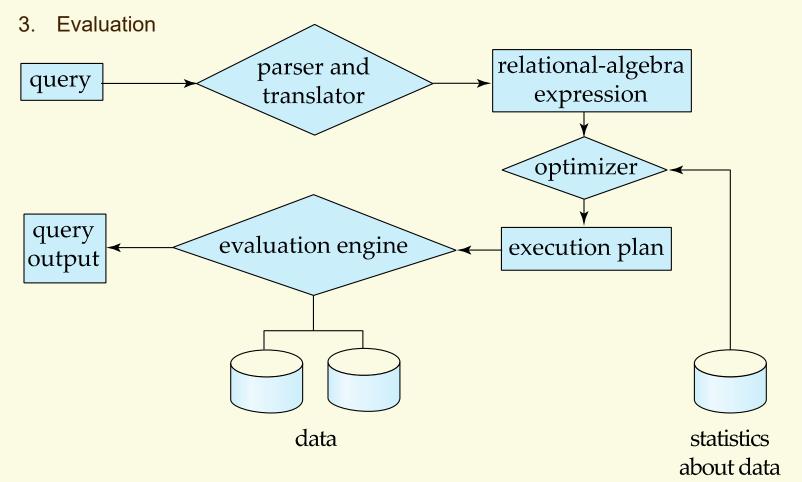
Query Processing

- ✓ Overview: query plan & optimization
- Measures of Query Cost
- ✓ Query optimization: Operators



Basic Steps in Query Processing

- 1. Parsing and translation
- 2. Optimization



Basic Steps in Query Processing

- Parsing and translation
 - translate the query into its internal form.
 - For RDBMS and SQL DB: relational algebra.
 - Parser checks syntax, verifies relations
- Optimization
 - Generate query(evaluation) plan from relational algebra
- ✓ Evaluation
 - The query-execution engine takes a query evaluation plan, executes that plan, and returns the answers to the query.

Basic Steps in Query Processing: Optimization

- A relational algebra expression may have many equivalent expressions
 - E.g., $\sigma_{salary<75000}(\Pi_{salary}(instructor))$ is equivalent to $\Pi_{salary}(\sigma_{salary<75000}(instructor))$
- Each relational algebra operation can be evaluated using one of several different algorithms
- ✓ Annotated expression specifying detailed evaluation strategy is called an evaluation-plan.
 - E.g., can use an index on salary to find instructors with salary <
 75000,
 - or can perform complete relation scan and discard instructors with salary ≥ 75000

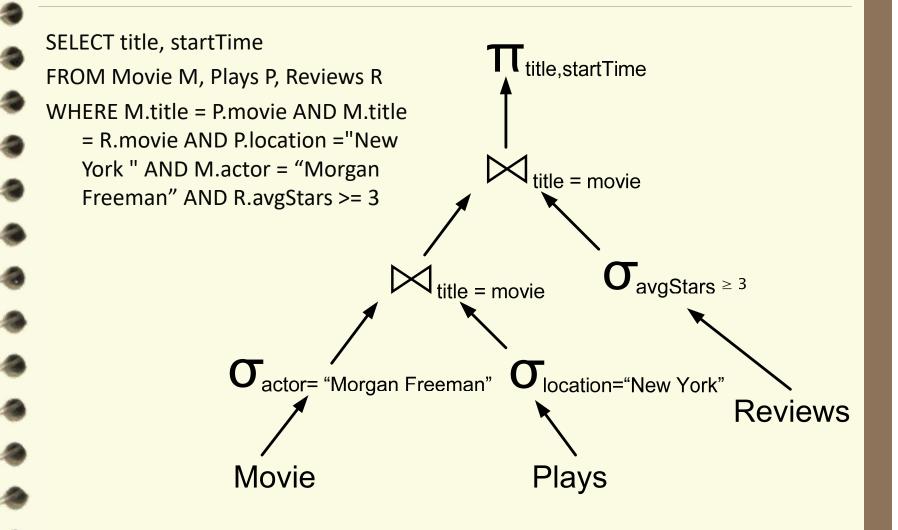
How to measure?

Example Query

SELECT title, startTime
FROM Movie M, Plays P, Reviews R
WHERE M.title = P.movie AND M.title
= R.movie AND P.location ="New
York " AND M.actor = "Morgan
Freeman" AND R.avgStars >= 3

 Query is parsed, generally broken into predicates, then converted into a logical query plan used by the optimizer

Example Logical Query Plan



Query Optimization

- ✓ Goal: compare all **equivalent** query expressions and their low-level implementations (**operators**)
- Can be divided into:
 - Plan enumeration ("search")
 - Cost estimation

Foundations of Query Plan Enumeration

Exploit properties of the relational algebra to find equivalent expressions, e.g.:

$$(R \bowtie S) \bowtie T = R \bowtie (S \bowtie T)$$

$$\sigma_{\alpha}(R \bowtie S) = (\sigma_{\alpha}(R) \bowtie S) \qquad \text{if } \alpha \text{ only refers to } R$$

- Assume selection, projection are always done as early as possible
- Joins (generally) satisfy principle of optimality:

Best way to compute RMSMT takes advantage of the best way of doing a 2 way join, followed by one additional join:

$$(R\bowtie S)\bowtie T$$
, $(R\bowtie T)\bowtie S$, or $(S\bowtie T)\bowtie R$

(uses optimal way of computing this expression)

Enumerating Plans

Can formulate as a dynamic programming problem:

- 1. Base case: consider all possible ways of accessing the base tables, with all selections & projections pushed down
- 2. Recursive case (i=2 ... number of joins+1): explore all ways to join results with (i-1) tables, with one additional table
 - Common heuristic only considers linear plans
- 3. Then repeat process for all Cartesian products
- 4. Apply **grouping** and aggregation

Find the Best Join Plan

```
procedure findbestplan(S) if (bestplan[S].cost \neq \infty) return bestplan[S]
```

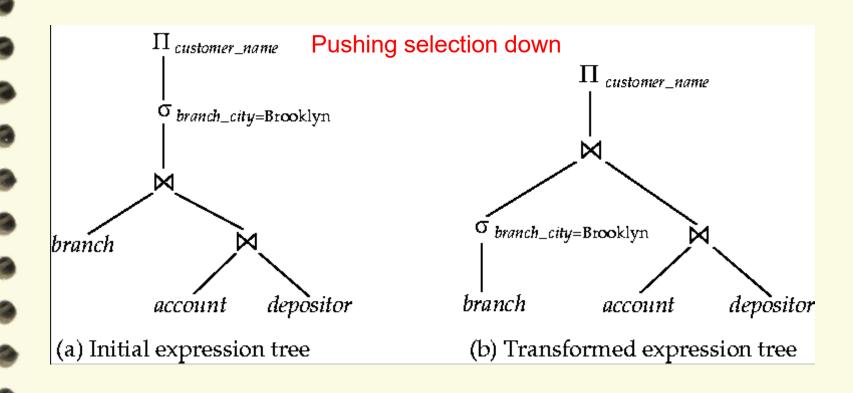
return bestplan[S]

Complexity: $O(3^n)$ Space complexity: $O(2^n)$

// else bestplan[S] has not been computed earlier, compute it now for each non-empty proper subset S1 of S
P1= findbestplan(S1) P2= findbestplan(S - S1)
A = best algorithm for joining results of P1 and P2 cost = P1.cost + P2.cost + cost of A
if cost < bestplan[S].cost
bestplan[S].cost = cost
bestplan[S].plan = "execute P1.plan; execute P2.plan;
join results of P1 and P2 using A"

Query Optimization Example

- Query: find the names of all customers who have an account at any branch located in Brooklyn
- Relational expression:



Enumeration of Equivalent Plans

•
$$\sigma_{\theta_1 \wedge \theta_2}(E) = \sigma_{\theta_1}(\sigma_{\theta_2}(E))$$

•
$$\sigma_{\theta_1}(\sigma_{\theta_2}(E)) = \sigma_{\theta_2}(\sigma_{\theta}(E))$$

•
$$\Pi_{L_{i}}(\Pi_{L_{i}}(K(\Pi_{L_{n}}(E))K)) = \Pi_{L_{i}}(E)$$

•
$$\sigma_{\theta}(E_1 \times E_2) = E_1 \bowtie_{\theta} E_2$$

•
$$\sigma_{\theta 1}(\mathsf{E}_1 \bowtie_{\theta 2} \mathsf{E}_2) = \mathsf{E}_1 \bowtie_{\theta 1 \land \theta 2} \mathsf{E}_2$$

•
$$E_1 \bowtie_{\theta} E_2 = E_2 \bowtie_{\theta} E_1$$

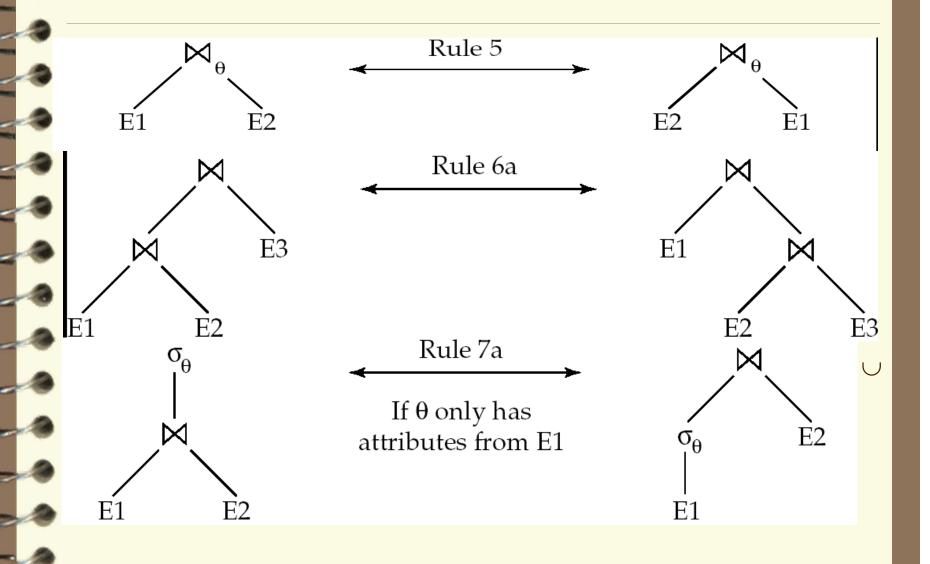
•
$$(E_1 \bowtie E_2) \bowtie E_3 = E_1 \bowtie (E_2 \bowtie E_3)$$

•
$$(E_1 \bowtie_{\theta_1} E_2) \bowtie_{\theta_2 \land \theta_3} E_3 = E_1 \bowtie_{\theta_2 \land \theta_3} (E_2 \bowtie_{\theta_2} E_3)$$

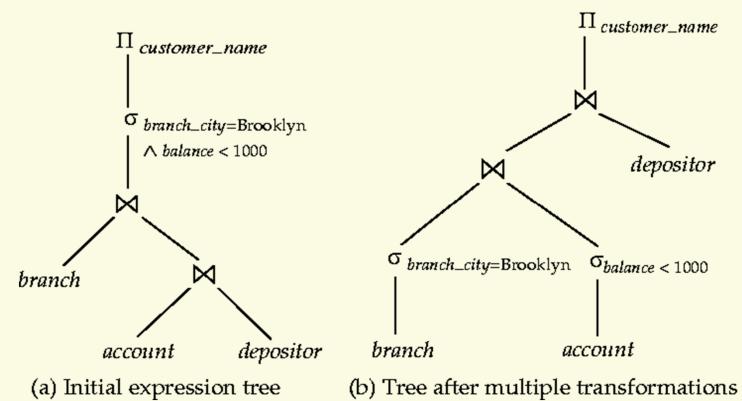
•
$$\sigma_{\theta 0}(\mathsf{E}_1 \bowtie_{\theta} \mathsf{E}_2) = (\sigma_{\theta 0}(\mathsf{E}_1)) \bowtie_{\theta} \mathsf{E}_2$$

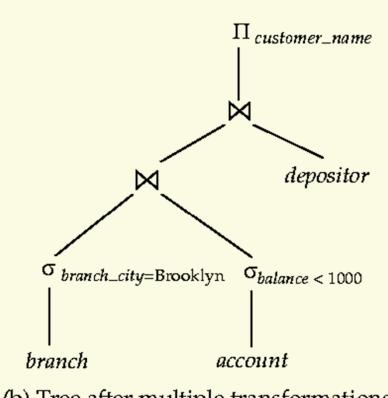
•
$$\sigma_{\theta_1} \wedge_{\theta_2} (\mathsf{E}_1 \bowtie_{\theta} \mathsf{E}_2) = (\sigma_{\theta_1}(\mathsf{E}_1)) \bowtie_{\theta} (\sigma_{\theta_2}(\mathsf{E}_2))$$

Enumeration of Equivalent Plans



Example







Measures of Query Cost

- Cost is generally measured as total elapsed time for answering query
 - Many factors contribute to time cost
 - disk accesses, CPU, or network communication
- Disk access is the predominant cost (I/O). Measured by
 - Number of seeks* average-seek-cost
 - Number of blocks read * average-block-read-cost
 - Number of blocks written * average-block-write-cost
 - Cost to write a block is greater than cost to read a block
 - data is read back after being written to ensure that the write was successful

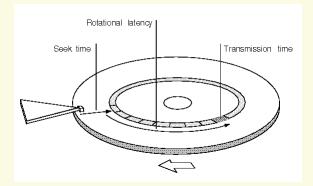
Measures of Query Cost (Cont.)

✓ I/O cost

- Use the number of block transfers from disk and the number
 of seeks as the cost measures
- t_T time to transfer one block
- $t_{\rm S}$ time for one seek
- Cost for b block transfers plus S seeks

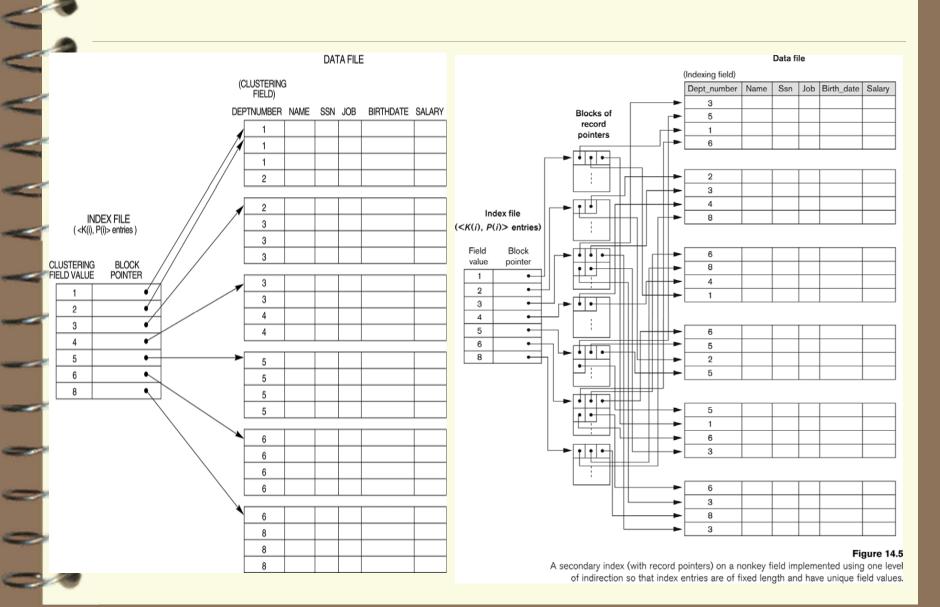
$$b * t_T + S * t_S$$

- ✓ CPU cost
 - for main-memory algorithms

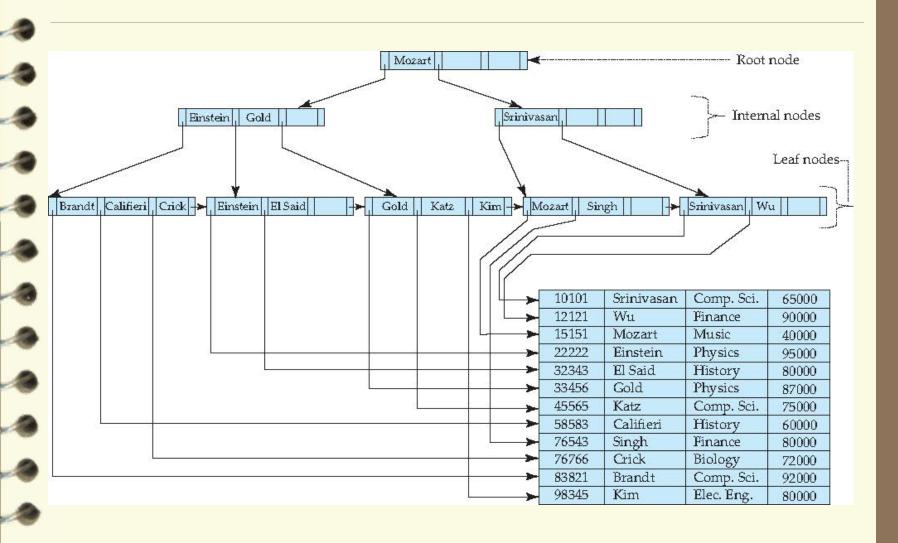


✓ Index are often used to reduce cost! – a flashback

Single-level index: Primary and Secondary



Multi-level index: Example of B+-Tree



Catalog information for cost estimation

- ✓ NR: # of tuples in a relation R
- ✓ BR: # of blocks that contain tuples of relation R
- ✓ SR: size of tuple of R
- ✓ FR: blocking factor; # of tuples from R that fit into one block:
 FR=NR/BR
- √ V(A, R): # of distinct value for attribute A in R.
- ✓ Sc(A,R): selectivity of attribute A = average number of tuples of R that satisfy an equality condition on A; Sc(A,R)=NR/V(A,R)

Query Plan Cost Estimation

For each expression, predict the **cost and output size** given what we know about the inputs

Requires significant information:

- A cost formula for each algorithm, in terms of disk
 I/Os, CPU speed, ...
- Calibration parameters for host machine performance
- Information about the **distributions** of join and grouping columns



Selections

- Selection: File scan
 - A1: Linear scan: Scan each file block and test. Cost: BR* t_T + t_S
- Selection: Index
 - A2 (primary index, equality on key). Retrieve a single record that satisfies the corresponding equality condition
 - $Cost = (h_i + 1) * (t_T + t_S)$ -- h_i : height of the (B+-tree)
 - A3 (primary index, equality on Nonkey) Retrieve multiple records.
 - Records will be on consecutive blocks
 - Let b = number of blocks containing matching records

•
$$Cost = h_i (t_T + t_S) + b^* t_T$$

= $h_i (t_T + t_S) + Sc(A,R)/FR * t_T$

NR: # of tuples

BR: # of blocks

SR: size of tuple of R

FR: block size

Sc(A,R): selectivity

Selections Using Indices

- A4 (secondary index, equality on nonkey).
 - Retrieve a single record if the search-key is a candidate key
 - $Cost = (h_i + 1) * (t_T + t_S)$
 - Retrieve multiple records if search-key is not a candidate key
 - each of n matching records may be on a different block
 - Cost = $(h_i + n) * (t_T + t_S)$
 - Can be very expensive!

NR: # of tuples

BR: # of blocks

SR: size of tuple of R

FR: block size Sc(A,R): selectivity

Selections Involving Comparisons

- ✓ Comparison: $\sigma_{A \le V}(r)$ or $\sigma_{A \ge V}(r)$
 - a linear file scan,
 - or by using indices in the following ways:
- ✓ A5 (primary index, comparison). (Relation is sorted on A)
 - For $\sigma_{A \ge V}(r)$ use index to find first tuple $\ge V$ and scan relation sequentially from there
 - For σ_{A≤V}(r) just scan relation sequentially till first tuple > v; do not use index
- ✓ A6 (secondary index, comparison).
 - scan index sequentially to find pointers to records
 - retrieve records that are pointed to
 - requires an I/O for each record
 - Linear file scan may be cheaper

NR: # of tuples

BR: # of blocks

SR: size of tuple of R

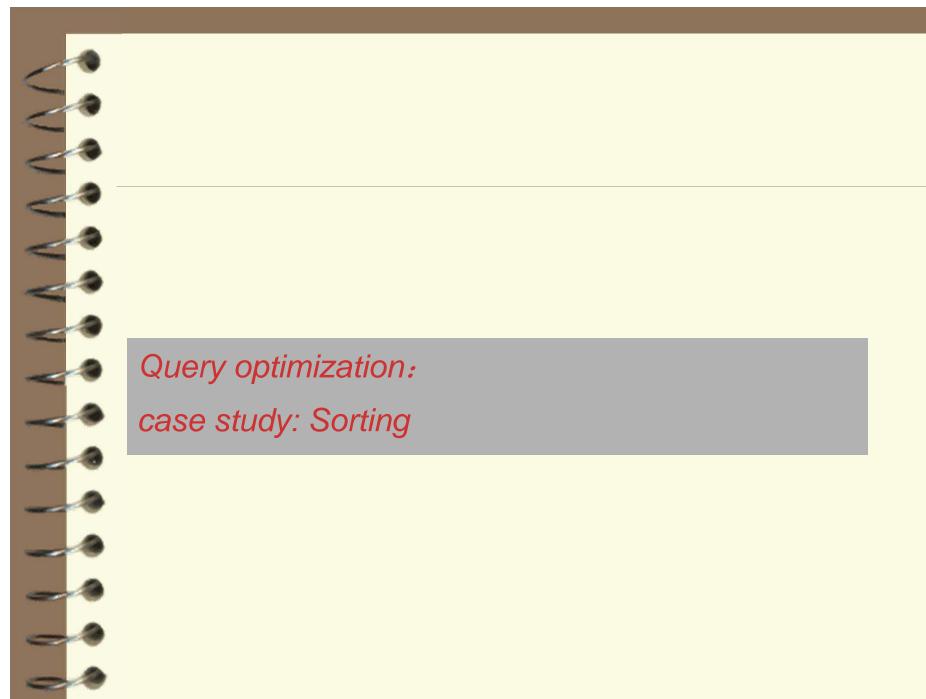
FR: block size Sc(A,R): selectivity

Implementation of Complex Selections

- ✓ Conjunction: $\sigma_{\theta 1} \land \theta_{2} \land \dots \theta_{n}(r)$
- ✓ A7 (conjunctive selection using one index).
 - Select a combination of θ_i that results in the least cost for $\sigma_{\theta_i}(r)$.
 - Test other conditions on tuple after fetching it into memory buffer.
- **✓ A8** (conjunctive selection using composite index).
 - Use appropriate composite (multiple-key) index if available.
- ✓ A9 (conjunctive selection by intersection of identifiers).
 - Requires indices with record pointers.
 - Use corresponding index for each condition, and take intersection of all the obtained sets of record pointers.
 - Then fetch records from file
 - If some conditions do not have appropriate indices, apply test in memory.

Algorithms for Complex Selections

- ✓ Disjunction: $\sigma_{\theta 1} \vee_{\theta 2} \vee \ldots_{\theta n} (r)$.
- A10 (disjunctive selection by union of identifiers).
 - Applicable if all conditions have available indices.
 - Otherwise use linear scan.
 - Use corresponding index for each condition, and take union of all the obtained sets of record pointers.
 - Then fetch records from file
- ✓ Negation: $\sigma_{-\theta}(r)$
 - Use linear scan on file
 - If very few records satisfy $\neg \theta$, and an index is applicable to θ
 - Find satisfying records using index and fetch from file



Sorting

- We may build an index on the relation, and then use the index to read the relation in sorted order. May lead to one disk block access for each tuple.
- ✓ For relations that fit in memory, techniques like quicksort can be used. For relations that don't fit in memory, external sort-merge is a good choice.

External Sort-Merge

Let *M* denote memory size (in pages).

1. Create sorted runs. Let i be 0 initially.

Repeatedly do the following till the end of the relation:

- (a) Read *M* blocks of relation into memory
- (b) Sort the in-memory blocks
- (c) Write sorted data to run file R_i ; increment i.

Let the final value of *i* be *N*

2. Merge the runs (next slide).....

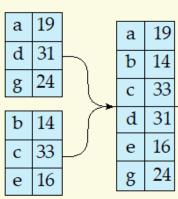
| g | 24 | |
|---|----|--|
| a | 19 | |
| d | 31 | |
| С | 33 | |
| b | 14 | |
| e | 16 | |

| a | 19 |
|---|----|
| d | 31 |
| g | 24 |
| | |
| b | 14 |
| С | 33 |
| е | 16 |

External Sort-Merge (Cont.)

- 2. Merge the runs (N-way merge). Assume that N < M.
 - 1. Use *N* blocks of memory to buffer input runs, and 1 block to buffer output. Read the first block of each run into its buffer page
 - 2. repeat
 - 1. Select the first record (in sort order) among all buffer pages
 - 2. Write the record to the output buffer. If the output buffer is full write it to disk.
 - Delete the record from its input buffer page.
 If the buffer page becomes empty then
 read the next block (if any) of the run into the buffer.
 - 3. until all input buffer pages are empty:

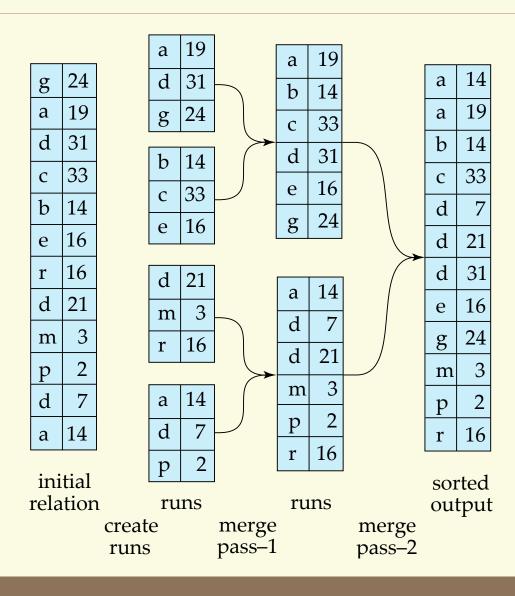
| g | 24 | |
|---|----|--|
| a | 19 | |
| d | 31 | |
| С | 33 | |
| b | 14 | |
| e | 16 | |



External Sort-Merge (Cont.)

- If $N \ge M$, several merge passes are required.
 - In each pass, contiguous groups of M 1 runs are merged.
 - A pass reduces the number of runs by a factor of *M*-1, and creates runs longer by the same factor.
 - E.g. If M=11, and there are 90 runs, one pass reduces the number of runs to 9, each 10 times the size of the initial runs
 - Repeated passes are performed till all runs have been merged into one.

Example: External Sorting Using Sort-Merge



External Merge Sort (Cont.)

Cost analysis:

- 1 block per run leads to too many seeks during merge
 - Instead use b_b buffer blocks per run
 - \rightarrow read/write b_b blocks at a time
 - Can merge \(\black M/b_b \) —1 runs in one pass
- Total number of merge passes required: $\lceil \log_{\lfloor M/bb \rfloor 1}(b_r/M) \rceil$.
- Block transfers for initial run creation as well as in each pass is 2b_r
 - for final pass, we don't count write cost
 - we ignore final write cost for all operations since the output of an operation may be sent to the parent operation without being written to disk
 - Thus total number of block transfers for external sorting: $b_r(2\lceil \log_{\lfloor M/bb \rfloor-1}(b_r/M) \rceil + 1)\lceil$

Seeks: next slide

NR: # of tuples

BR: # of blocks

SR: size of tuple of R

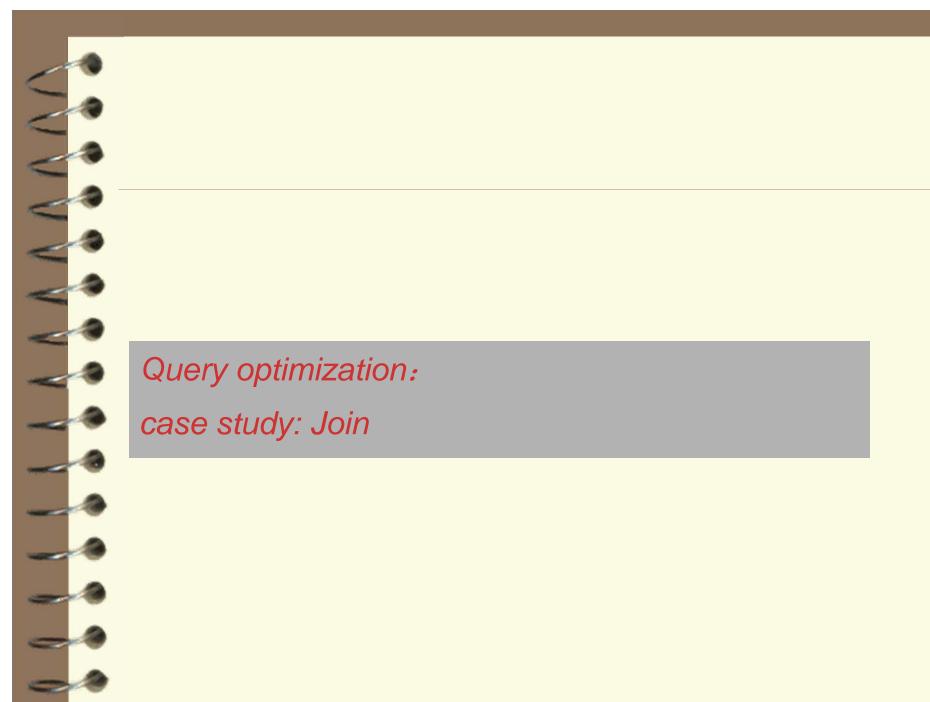
FR: block size

Sc(A,R): selectivity

External Merge Sort (Cont.)

- Cost of seeks
 - During run generation: one seek to read each run and one seek to write each run
 - $2\lceil b_r/M \rceil$
 - During the merge phase
 - Need $2 \lceil b_r / b_b \rceil$ seeks for each merge pass
 - except the final one which does not require a write
 - Total number of seeks:

$$2\lceil b_r/M \rceil + \lceil b_r/b_b \rceil (2\lceil \log_{|M/bb|-1}(b_r/M) \rceil - 1)$$



Join Operation

- ✓ Several different algorithms to implement joins
 - Nested-loop join
 - Block nested-loop join
 - Indexed nested-loop join
 - Merge-join
 - Hash-join
- Choice based on cost estimate
- Examples use the following information
 - Number of records of student: 5,000 takes: 10,000
 - Number of blocks of student: 100 takes: 400

Nested-Loop Join

To compute the theta join $r \bowtie_{\Theta} s$ for each tuple t_r in r do begin for each tuple t_s in s do begin test pair (t_r, t_s) to see if they satisfy the join condition θ if they do, add $t_r \bullet t_s$ to the result.

- end
- \checkmark r is called the **outer relation** and s the **inner relation** of the join.
- Requires no indices and can be used with any kind of join condition.
- Expensive since it examines every pair of tuples in the two relations.

Nested-Loop Join (Cont.)

In the worst case, if enough memory only to hold one block of each relation, the estimated cost is

$$n_r * b_s + b_r$$
 block transfers, plus $n_r + b_r$ seeks

If the smaller relation fits entirely in memory, use that as the inner relation.

- Reduces cost to $b_r + b_s$ block transfers and 2 seeks

Assuming worst case memory availability cost estimate is

with student as outer relation:

5000 * 400 + 100 = 2,000,100 block transfers,

• 5000 + 100 = 5100 seeks

with takes as the outer relation

• 10000 * 100 + 400 = 1,000,400 block transfers and 10,400 seeks

If smaller relation (*student*) fits entirely in memory, the cost estimate will be 500 block transfers.

Block nested-loops algorithm (next slide) is preferable.

Block Nested-Loop Join

Variant of nested-loop join in which every block of inner relation is paired with every block of outer relation.

```
for each block B_r of r do begin
   for each block B_s of s do begin
       for each tuple t_r in B_r do begin
           for each tuple t_s in B_s do begin
              Check if (t_r, t_s) satisfy the join condition
              if they do, add t_r \cdot t_s to the result.
           end
       end
   end
end
```

Block Nested-Loop Join (Cont.)

Worst case estimate: $b_r * b_s + b_r$ block transfers + 2 * b_r seeks

Each block in the inner relation s is read once for each block in the outer relation

Best case: $b_r + b_s$ block transfers + 2 seeks.

Use index on inner relation if available (next slide)

Indexed Nested-Loop Join

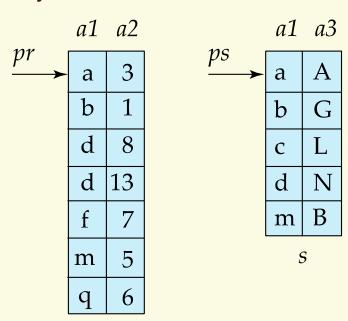
- For each tuple t_r in the outer relation r, use the index to look up tuples in s that satisfy the join condition with tuple t_r .
- ✓ Worst case: buffer has space for only one page of r, and, for each tuple in r, we perform an index lookup on s.
- \checkmark Cost of the join: $b_r(t_T + t_S) + n_r * c$
 - Where c is the cost of traversing index and fetching all matching s tuples for one tuple or r
 - c can be estimated as cost of a single selection on s using the join condition.
- If indices are available on join attributes of both *r* and *s*, use the relation with fewer tuples as the outer relation.

Example of Nested-Loop Join Costs

- \checkmark Compute *student* \bowtie *takes,* with *student* as the outer relation.
- ✓ Let *takes* have a primary B+-tree index on the attribute *ID*, which contains 20 entries in each index node.
- ✓ Since *takes* has 10,000 tuples, the height of the tree is 4, and one more access is needed to find the actual data
- √ student has 5000 tuples
- Cost of block nested loops join
 - -400*100 + 100 = 40,100 block transfers + 2 * 100 = 200 seeks
 - assuming worst case memory
 - may be significantly less with more memory
- ✓ Cost of indexed nested loops join
 - 100 + 5000 * 5 = 25,100 block transfers and seeks.
 - CPU cost likely to be less than that for block nested loops join

Merge-Join

- 1. Sort both relations on their join attribute (if not already sorted on the join attributes).
- 2. Merge the sorted relations to join them
 - 1. Join step is similar to the merge stage of the sort-merge algorithm.
 - 2. Main difference is handling of duplicate values in join attribute every pair with same value on join attribute must be matched



Merge-Join (Cont.)

- Can be used only for equi-joins and natural joins
- ✓ Each block needs to be read only once (assuming all tuples for any given value of the join attributes fit in memory
- ✓ Thus the cost of merge join is:

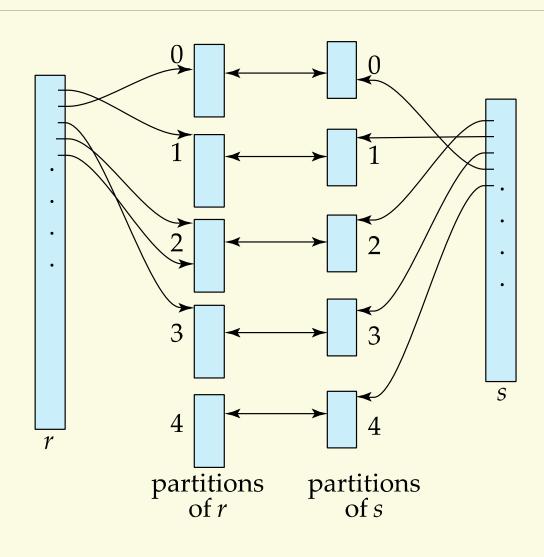
$$b_r + b_s$$
 block transfers $+ \lceil b_r / b_b \rceil + \lceil b_s / b_b \rceil$ seeks

- + the cost of sorting if relations are unsorted.
- ✓ hybrid merge-join: If one relation is sorted, and the other has a secondary B⁺-tree index on the join attribute
 - Merge the sorted relation with the leaf entries of the B⁺-tree .
 - Sort the result on the addresses of the unsorted relation's tuples
 - Scan the unsorted relation in physical address order and merge with previous result, to replace addresses by the actual tuples
 - Sequential scan more efficient than random lookup

Hash-Join

- Applicable for equi-joins and natural joins.
- ✓ A hash function *h* is used to partition tuples of both relations
- ✓ *h* maps *JoinAttrs* values to {0, 1, ..., *n*}, where *JoinAttrs* denotes the common attributes of *r* and *s* used in the natural join.
 - r_0, r_1, \ldots, r_n denote partitions of r tuples
 - Each tuple $t_r \in r$ is put in partition r_i where $i = h(t_r [JoinAttrs])$.
 - $-s_0$, s_1 ..., s_n denotes partitions of s tuples
 - Each tuple $t_S \in s$ is put in partition s_i , where $i = h(t_S = [JoinAttrs])$.

Hash-Join (Cont.)



Hash-Join (Cont.)

- \checkmark r tuples in r_i need only to be compared with s tuples in s_i Need not be compared with s tuples in any other partition, since:
 - an r tuple and an s tuple that satisfy the join condition
 will have the same value for the join attributes.
 - If that value is hashed to some value i, the r tuple has to be in r_i and the s tuple in s_i .

Hash-Join Algorithm

The hash-join of *r* and *s* is computed as follows.

- 1. Partition the relation *s* using hashing function *h*. When partitioning a relation, one block of memory is reserved as the output buffer for each partition.
- 2. Partition *r* similarly.
- 3. For each i:
 - (a) Load s_i into memory and build an in-memory hash index on it using the join attribute. This hash index uses a different hash function than the earlier one h.
 - (b) Read the tuples in r_i from the disk one by one. For each tuple t_r locate each matching tuple t_s in s_i using the inmemory hash index. Output the concatenation of their attributes.

Relation s is called the **build input** and r is called the **probe input**.

Hash-Join algorithm (Cont.)

- The # of partitions n and the hash function h is chosen such that each s_i should fit in memory.
 - Typically n is chosen as \[\bar{b}_s/M \] * f where f is a "fudge factor", typically around 1.2
 - The probe relation partitions r_i need not fit in memory
- ✓ Recursive partitioning required if number of partitions n is greater than number of pages M of memory.
 - instead of partitioning n ways, use M-1 partitions for s
 - Further partition the M-1 partitions using a different hash function
 - Use same partitioning method on r
 - Rarely required: e.g., with block size of 4 KB, recursive partitioning not needed for relations of < 1GB with memory size of 2MB, or relations of < 36 GB with memory of 12 MB

Cost of Hash-Join

✓ If recursive partitioning is not required: cost of hash join is

$$3(b_r + b_s) + 4 * n_h$$
 block transfers +

$$2(\lceil b_r/b_b \rceil + \lceil b_s/b_b \rceil)$$
 seeks

- ✓ If recursive partitioning required:
 - number of passes required for partitioning build relation s to less than M blocks per partition is $\lceil log_{M/bb} \mid -1 (b_s/M) \rceil$
 - best to choose the smaller relation as the build relation.
 - Total cost estimate is:

$$2(b_r + b_s) \lceil log_{M/bb - 1}(b_s/M) \rceil + b_r + b_s$$
 block transfers + $2(\lceil b_r/b_b \rceil + \lceil b_s/b_b \rceil) \lceil log_{M/bb - 1}(b_s/M) \rceil$ seeks

- ✓ If the entire build input can be kept in main memory no partitioning is required
 - Cost estimate goes down to $b_r + b_s$.

Example of Cost of Hash-Join

instructor ⋈ *teaches*

- Assume that memory size is 20 blocks
- ✓ $b_{instructor}$ = 100 and $b_{teaches}$ = 400.
- ✓ instructor is to be used as build input. Partition it into five partitions, each of size 20 blocks. This partitioning can be done in one pass.
- ✓ Similarly, partition *teaches* into five partitions,each of size 80. This is also done in one pass.
- ✓ Therefore total cost, ignoring cost of writing partially filled blocks:
 - 3(100 + 400) = 1500 block transfers + $2(\lceil 100/3 \rceil + \lceil 400/3 \rceil) = 336$ seeks

Complex Joins

✓ Join with a conjunctive condition:

$$r \bowtie_{\theta 1 \land \theta 2 \land \dots \land \theta n} s$$

- Either use nested loops/block nested loops, or
- Compute the result of one of the simpler joins $r \bowtie_{\theta i} s$
 - final result comprises those tuples in the intermediate result that satisfy the remaining conditions

$$\theta_1 \wedge \ldots \wedge \theta_{i-1} \wedge \theta_{i+1} \wedge \ldots \wedge \theta_n$$

✓ Join with a disjunctive condition

$$r \bowtie_{\theta 1 \vee \theta 2 \vee \dots \vee \theta n} s$$

- Either use nested loops/block nested loops, or
- Compute as the union of the records in individual joins $r \bowtie_{\theta} s$:

$$(r \bowtie_{\theta_1} s) \cup (r \bowtie_{\theta_2} s) \cup \ldots \cup (r \bowtie_{\theta_n} s)$$

"Big O" notation

Given two functions, f and g, say that "f is of order g" if

- there is a constant c, and
- a value x_0

such that

Apart from a fixed multiplicative constant, the function g is an

- upper bound on the function f
- valid for large values of its argument.

Notation: write to mean "f is of order g".

Sometimes write to remind us the arguments.

Example: N-by-N matrix, N-by-1 vector, multiply

```
Y = zeros(N,1); initialize space, c<sub>1</sub>N

for i=1:N initialize "for" loop, c<sub>2</sub>N

Y(i) = 0.0; Scalar assignment, c<sub>3</sub>

for j=1:N initialize "for" loop, c<sub>2</sub>N

Y(i) = Y(i) + A(i,j)*x(j); C<sub>4</sub>

end End of loop, return/exit, c<sub>5</sub>

End of loop, return/exit, c<sub>5</sub>
```

$$Total = c_1 N + c_2 N + N(c_3 + c_2 N + N(c_4 + c_5) + c_5)$$

$$= (c_2 + c_4 + c_5)N^2 + (c_1 + c_2 + c_3 + c_5)N$$

$$= c_6 N^2 + c_7 N$$