



CPT-S 415

Big Data

Yinghui Wu

EME B45

CPT-S 415

Big Data

Graph Query Processing

- ✓ Basics of Graph Algorithms
 - Graph search (traversal)
 - PageRank
 - Nearest neighbors
 - Keyword search
 - Graph pattern matching (next lecture)

Graph query processing: operators and algorithms

When it comes to graphs ...

- ✓ Semistructured:
 - No schema
 - No constraints yet
- ✓ No standard query languages
 - A variety of queries used in practice
 - Nontrivial
- ✓ What is the complexity of the following problems?
 - Subgraph isomorphism NP-complete
 - Simple path: given a graph G , a pair (s, t) of nodes in G , and a regular expression R , it is to decide whether there exists a simple path from s to t that satisfies R .
- ✓ Query optimization techniques, indexing, updates, ... preliminary

The study of graph queries is still in its infancy

Basic graph queries and algorithms

- ✓ Graph search (traversal)
- ✓ PageRank
- ✓ Nearest neighbors
- ✓ Keyword search
- ✓ Graph pattern matching (a full treatment of itself)

Widely used in graph algorithms

Graph search (traversal)

Path queries

✓ Reachability

- Input: A directed graph G , and a pair of nodes s and t in G
- Question: Does there exist a path from s to t in G ?

✓ Distance

- Input: A directed weighted graph G , and a node s in G
- Output: The lengths of **shortest paths** from s to all nodes in G

✓ Regular path

- Input: A node-labeled directed graph G , a pair of nodes s and t in G , and a regular expression R
- Question: Does there exist a (simple) path p from s to t that satisfies R ?

Reachability queries

✓ Reachability

- Input: A directed graph G , and a pair of nodes s and t in G
- Question: Does there exist a path from s to t in G ?

✓ Applications: a routine operation

- Social graphs: are two people related for security reasons?
- Biological networks: find genes that are (directly or indirectly) influenced by a given molecule

Nodes: molecules, reactions or physical interactions

Edges: interactions

How to evaluate reachability queries?

Breadth-first search

✓ BFS (G, s, t):


1. while Que is nonempty do
 - a. $v \leftarrow \text{Que.dequeue}()$;
 - b. if $v = t$ then return true;
 - c. for all adjacent edges $e = (v, u)$ of v do
 - a) if not flag(u)
then flag(u) \leftarrow true; enqueue u onto Que;
2. return false

Use (a) a queue Que, initialized with s, (b) flag(v) for each node, initially false; and (c) adjacency lists to store G

Breadth-first, by using a queue

Why do we need the test?

- ✓ Complexity: each node and edge is examined at most once



What is the

Breadth-first search

Too costly as a routine operation
when G is large

✓ Reachability: **NL-complete**

✓ **BFS** (G, s, t): $O(|V| + |E|)$ time and space

1. while Que is nonempty do
 - a. $v \leftarrow \text{Que.dequeue}()$;
 - b. if $v = t$ then return true;
 - c. for all adjacent edges $e = (v, u)$ of v do
 - a) if not flag(u)
then flag(u) \leftarrow true; enqueue u onto Que;
2. return false

✓ **$O(1)$ time?** Yes, adjacency matrix, but **$O(|V|^2)$ space**

How to trike a balance?

2-hop cover

- ✓ For each node v in G ,

$$2hop(v) = (L_{in}(v), L_{out}(v))$$

- $L_{in}(v)$: a set of nodes in G that can reach v
- $L_{out}(v)$: a set of nodes in G that v can reach

- ✓ To ensure: node s can reach t if and only if

$$L_{out}(s) \cap L_{in}(t) \neq \emptyset$$


- Testing: better than $O(|V| + |E|)$ on average
- Space: $O(|V| |E|^{1/2})$

Find a minimum 2-hop cover? NP-hard

Maintenance cost in response to changes to G

A number of algorithms for reachability queries (see reading list)

Distance queries



Edsger Wybe Dijkstra
(1930-2002)

- ✓ Distance: single-source shortest-path problem
 - Input: A directed weighted graph G , and a node s in G
 - Output: The lengths of **shortest paths** from s to all nodes in G
- ✓ Application: transportation networks

✓ Dijkstra (G, s, w):

1. for all nodes v in V do

a. $d[v] \leftarrow \infty$; \leftarrow

2. $d[s] \leftarrow 0$; $Que \leftarrow V$;

3. while Que is nonempty do

a. $u \leftarrow \text{ExtractMin}(Que)$;


b. for all nodes v in $\text{adj}(u)$ do

a) if $d[v] > d[u] + w(u, v)$ then $d[v] \leftarrow d[u] + w(u, v)$;

Use a priority queue Que ; $w(u, v)$: weight of edge (u, v) ; $d(u)$: the distance from s to u


Extract one with the minimum $d(u)$

Example: Dijkstra's algorithm



$Q = \{s, a, b, c, d\}$
 $d: \{(a, \infty), (b, \infty), (c, \infty), (d, \infty)\}$


Example: Dijkstra's algorithm



$Q = \{a, b, c, d\}$


$d: \{(a, 10), (b, \infty), (c, 5), (d, \infty)\}$

Example: Dijkstra's algorithm



$Q = \{a, b, d\}$
 $d: \{(a, 8), (b, 14), (c, 5), (d, 7)\}$


Example: Dijkstra's algorithm



$Q = \{a, b\}$


$d: \{(a, 8), (b, 13), (c, 5), (d, 7)\}$

Example: Dijkstra's algorithm



$Q = \{b\}$
 $d: \{(a, 8), (b, 9), (c, 5), (d, 7)\}$

Example: Dijkstra's algorithm



$Q = \{\}$

$d: \{(a,8), (b,9), (c,5), (d,7)\}$


How to speed it up?

$O(|V| \log |V| + |E|)$. A beaten-to-death topic ?

Landmarks

✓ Landmark vectors

- A list of nodes L in a graph G , s.t for each pair (u,v) of nodes in G , there is an node in L on a shortest path from u to v
- Answering distance query: linear time



A landmark vector LM

Regular path queries

✓ Regular simple path

- Input: A node-labeled directed graph G , a pair of nodes s and t in G , and a regular expression R
- Question: Does there exist a simple path p from s to t such that the labels of p match R ?

A. Mendelzon, and P. T. Wood. *Finding Regular Simple Paths In Graph Databases*.
SICOMP 24(6), 1995

- ✓ NP-complete, even when R is a fixed regular expression $(00)^*$ or 0^*10^* .
- ✓ In PTIME when G is a DAG (directed acyclic graph)

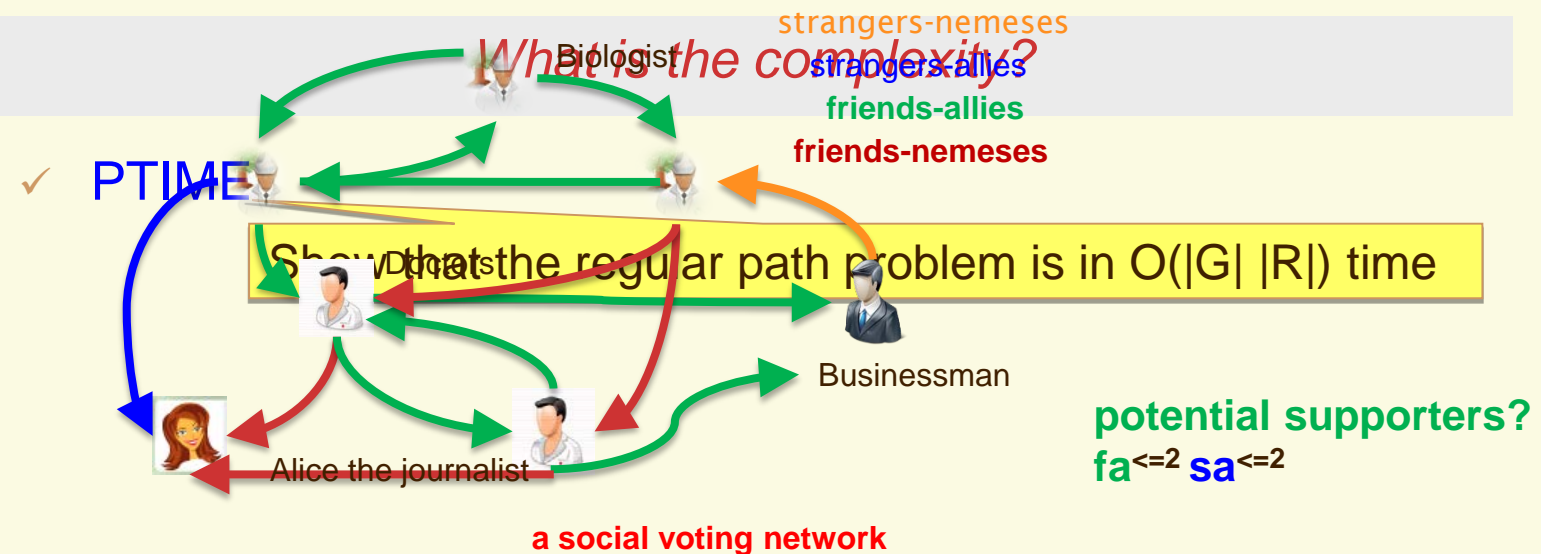
Patterns of social links

Why do we care about regular path queries?

Regular path queries

✓ Regular path

- Input: A node-labeled directed graph G , a pair of nodes s and t in G , and a **regular expression** R
- Question: Does there exist **a path** p from s to t such that the labels of adjacent nodes on p form a string in R ?



Graph queries are nontrivial, even for path queries

Strongly connected components

- ✓ A strongly connected component in a direct graph G is a set V of nodes in G such that
 - for any pair (u, v) in V , u can reach v and v can reach u ; and
 - V is **maximal**: adding any node to V makes it no longer strongly connected

- ✓ **SCC**

Find social circles: how large? How many?

- Input: A graph G
- Question: all strongly connected components of G

What is the complexity?

by extending search algorithms, e.g., BFS

$O(|V| + |E|)$

A spiral-bound notebook with a cream-colored page and a brown cover. The spiral binding is on the left side. A horizontal line is drawn across the page, and a gray rectangular box is positioned below it.

PageRank

Introduction to PageRank

✓ To measure the “quality” of a Web page

- Input: A directed graph G modelling the Web, in which nodes represent Web pages, and edges indicate hyperlinks
- Output: For each node v in G , $P(v)$: the likelihood that a random walk over G will arrive at v .

✓ Intuition: how a random walk can reach v ?

- A random jump: $\alpha (1/|V|)$ The chances of hitting v among $|V|$ pages
 - α : random jump factor (teleportation factor)
- Following a hyperlink: $(1 - \alpha) \sum_{(u \in L(v))} P(u)/C(u)$
 - $(1 - \alpha)$: damping factor
 - $(1 - \alpha) \sum_{(u \in L(v))} P(u)/C(u)$: the chances for one to click a hyperlink at a page u and reach v

Intuition

- ✓ Following a hyperlink: $(1 - \alpha) \sum_{(u \in L(v))} P(u)/C(u)$
 - $L(v)$: the set of pages that **link to** v ;
 - $C(u)$: the out-degree of node u (the number of links on u)
 - $P(u)/C(u)$:

The chances of reaching page v from page u

 - the probability of u being visited itself
 - the probability of clicking the link to v among $C(u)$ many links on page u
- ✓ Intuition:
 - the more pages link to v , and
 - the more popular those pages that link to v , v has a higher chance to be visited

One of the models

Putting together

The likelihood that page v is visited by a random walk:

$$\alpha (1/|V|) + (1 - \alpha) \sum_{u \in L(v)} P(u)/C(u)$$

random jump

following a link from other pages

✓ Recursive computation: for each page v in G ,

- compute $P(v)$ by using $P(u)$ for all $u \in L(v)$

until

too expensive; use an error factor


- converge: no changes to any $P(v)$
- after a fixed number of iterations

costly: trillions of pages

Parallel computation

How to speed it up?

An example of Simplified PageRank




$$M = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1 \\ 0 & 1/2 & 0 \end{bmatrix}$$

$$\begin{bmatrix} \text{yahoo} \\ \text{Amazon} \\ \text{Microsoft} \end{bmatrix} = \begin{bmatrix} 1/3 \\ 1/3 \\ 1/3 \end{bmatrix}$$

$$\begin{bmatrix} 1/3 \\ 1/2 \\ 1/6 \end{bmatrix} = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1 \\ 0 & 1/2 & 0 \end{bmatrix} \begin{bmatrix} 1/3 \\ 1/3 \\ 1/3 \end{bmatrix}$$

PageRank Calculation: first iteration

An example of Simplified PageRank




$$M = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1 \\ 0 & 1/2 & 0 \end{bmatrix}$$

$$\begin{bmatrix} \text{yahoo} \\ \text{Amazon} \\ \text{Microsoft} \end{bmatrix} = \begin{bmatrix} 1/3 \\ 1/3 \\ 1/3 \end{bmatrix}$$

$$\begin{bmatrix} 5/12 \\ 1/3 \\ 1/4 \end{bmatrix} = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1 \\ 0 & 1/2 & 0 \end{bmatrix} \begin{bmatrix} 1/3 \\ 1/2 \\ 1/6 \end{bmatrix}$$

PageRank Calculation: second iteration

An example of Simplified PageRank



$$M = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1 \\ 0 & 1/2 & 0 \end{bmatrix}$$

$$\begin{bmatrix} \text{yahoo} \\ \text{Amazon} \\ \text{Microsoft} \end{bmatrix} = \begin{bmatrix} 1/3 \\ 1/3 \\ 1/3 \end{bmatrix}$$

$$\begin{bmatrix} 3/8 \\ 11/24 \\ 1/6 \end{bmatrix} \quad \begin{bmatrix} 5/12 \\ 17/48 \\ 11/48 \end{bmatrix} \quad \dots \quad \begin{bmatrix} 2/5 \\ 2/5 \\ 1/5 \end{bmatrix}$$

Convergence after some iterations

Nearest neighbors

Nearest neighbor

✓ Nearest neighbor (kNN)

- Input: A set S of points in a space M , a query point p in M , a distance function $\text{dist}(u, v)$, and a positive integer k
- Output: Find top- k points in S that are closest to p based on $\text{dist}(p, u)$

Euclidean distance, Hamming distance, continuous variables, ...

✓ Applications

- POI recommendation: find me top- k restaurants close to where I am
- Classification: classify an object based on its nearest neighbors
- Regression: property value as the average of the values of its k nearest neighbors

Linear search, space partitioning, locality sensitive hashing, compression/clustering based search, ...

A number of techniques

kNN join

✓ kNN join

- Input: Two datasets R and S , a distance function $\text{dist}(r, s)$, and a positive integer k
- Output: pairs (r, s) for all r in R , where s is in S , and is one of the k -nearest neighbors of r

✓ A naive algorithm




Pairwise comparison

- Scanning S once for each object in R
- $O(|R| |S|)$: expensive when R or S is large

Can we do better?

Blocking and windowing





✓ blocking



only pairs in the same block are compared

GORDER: An Efficient Method for KNN Join Processing. VLDB 2004.

✓ windowing



window of a fixed size; only pairs in the same window are compared;

Several indexing and ordering techniques

A spiral-bound notebook with a cream-colored page and a brown cover. The spiral binding is on the left side. A horizontal line is drawn across the page. A grey rectangular box is positioned in the middle of the page, containing the text "Keyword search" in red, italicized font.

Keyword search

Keyword search

- ✓ Input: A list Q of keywords, a graph G , a positive integer k
- ✓ Output: top- k “matches” of Q in G

Information retrieval



Query Q : [*Jaguar*, *America*, *history*]



- ✓ What makes a match?
- ✓ How to sort the matches?
- ✓ How to efficiently find top- k matches?

Questions to answer

Semantics: Steiner tree

✓ Input: A list Q of keywords, a graph G , a weight function $w(e)$ on the edges on G , and a positive integer k

✓ Output: top- k Steiner trees that match Q

PageRank scores

✓ Match: a subtree T of G such that

- each keyword in Q is contained in a leaf of T

✓ Ranking:

- The total weight of T (the sum of $w(e)$ for all edges e in T)

The cost to connect the keywords

✓ Complexity?

NP-complete

What can we do about it?

Semantics: distinct-root (tree)

- ✓ Input: A list Q of keywords, a graph G , and a positive integer k
- ✓ Output: top- k distinct trees that match Q
- ✓ Match: a subtree T of G such that
 - each keyword in Q is contained in a leaf of T
- ✓ Ranking:
 - $\text{dist}(r, q)$: from the root of T to a leaf q
 - The sum of distances from the root to all leaves of T
- ✓ Diversification:
 - each match in the top- k answer has a distinct root

$$O(|Q| (|V| \log |V| + |E|))$$

Semantics: Steiner graphs

- ✓ Input: A list Q of keywords, an **undirected** (unweighted) graph G , a **positive integer r** , and a positive integer k
- ✓ Output: Find all r -radius Steiner graphs that match Q
- ✓ Match: a subgraph G' of G such that it is
 - **r -radius**: the shortest distance between any pair of nodes in G is at most r (at least one pair with the distance); and
 - **each keyword** is contained in either
 - a content node: containing the key word
 - a Steiner node: on a simple path between a pair of content nodes
- ✓ Computation: M^r , the r -th power of adjacency graph of G

Revision: minimum subgraphs

Answering keyword queries

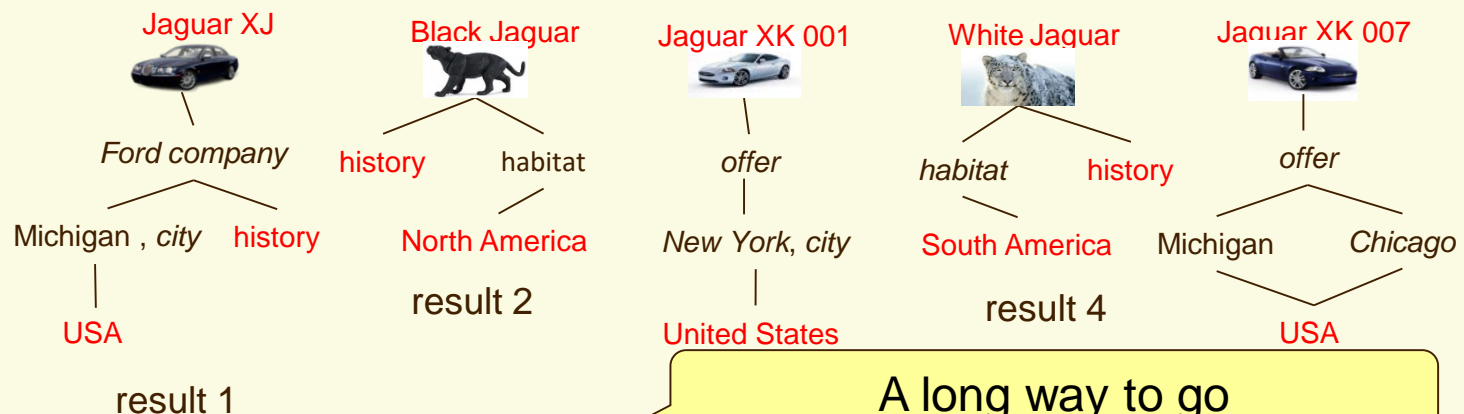

- ✓ A host of techniques
 - Backward search
 - Bidirectional search
 - Bi-level indexing
 - ...
- ✓ G. Bhalotia, A. Hulgeri, C. Nakhe, S. Chakrabarti, and S. Sudarshan. Keyword searching and browsing in databases using BANKS. ICDE 2002.
- ✓ V. Kacholia, S. Pandit, S. Chakrabarti, S. Sudarshan, R. Desai, and H. Karambelkar. Bidirectional expansion for keyword search on graph databases. VLDB 2005.
- ✓ H. He, H. Wang, J. Yang, and P. S. Yu. BLINKS: ranked keyword searches on graphs. SIGMOD 2007.

A well studied topic

However, ...

What does the user really want to find? Tree or graph? How to explain matches found?

✓ The semantics is rather “ad hoc”
Query Q: [Jaguar, 'America', 'history']



Add semantics to keyword search

A spiral-bound notebook with a cream-colored page and a brown cover. The spiral binding is on the left side. A horizontal line is drawn across the page. A grey rectangular box is positioned in the middle of the page, containing the text "Summing up".

Summing up

Reading List 2: Graph query languages

There have been efforts to develop

Querying topological structures

- ✓ SoQL: an SQL-like language to retrieve paths
- ✓ CRPQ: extending conjunctive queries with regular path expressions
 - R. Ronen and O. Shmueli. SoQL: A language for querying and creating data in social networks. ICDE, 2009.
 - P. Barceló, C. A. Hurtado, L. Libkin, and P. T. Wood. Expressive languages for path queries over graph-structured data. In PODS, 2010
- ✓ SPARQL: for RDF data
 - <http://www.w3.org/TR/rdf-sparql-query/>

Read this

Unfortunately, no “standard” query language for graphs, yet

Summary and review

- ✓ Review of searching RDBMS
 - Algorithms for selection/complex selection
 - Algorithms for joins: what are several common ways to compute joins? Design ideas?
- ✓ Why are reachability queries? Regular path queries? Connected components? Complexity? Algorithms?
- ✓ What are factors for PageRank? How does PageRank work?
- ✓ What are kNN queries? What is a kNN join? Complexity?
- ✓ What are keyword queries? What is its Steiner tree semantics? Distinct-root semantics? Steiner graph semantics? Complexity?
- ✓ Name a few applications of graph queries you have learned.
- ✓ Find graph queries that are not covered in the lecture.

Reading List 2 (graduate students)

- ❖ G. Bhalotia, A. Hulgeri, C. Nakhe, S. Chakrabarti, and S. Sudarshan. Keyword searching and browsing in databases using BANKS. ICDE 2002.
- ❖ V. Kacholia, S. Pandit, S. Chakrabarti, S. Sudarshan, R. Desai, and H. Karambelkar. Bidirectional expansion for keyword search on graph databases. VLDB 2005.
- ❖ H. He, H. Wang, J. Yang, and P. S. Yu. BLINKS: ranked keyword searches on graphs. SIGMOD 2007.
- ❖ W. Fan, J. Li, S. Ma, N. Tang, and Y. Wu. Adding Regular Expressions to Graph Reachability and Pattern Queries, ICDE 2011.
- ❖ M. R. Henzinger, T. Henzinger, and P. Kopke. Computing simulations on finite and infinite graphs. FOCS, 1995. <http://infoscience.epfl.ch/record/99332/files/HenzingerHK95.pdf>
- ❖ L. P. Cordella, P. Foggia, C. Sansone, M. Vento. A (Sub)Graph Isomorphism Algorithm for Matching Large Graphs, IEEE Trans. Pattern Anal. Mach. Intell. 26, 2004 (search Google scholar)
- ❖ A. Fard, M. U. Nisar, J. A. Miller, L. Ramaswamy, Distributed and scalable graph pattern matching: models and algorithms. Int. J. Big Data. http://cobweb.cs.uga.edu/~ar/papers/IJBD_final.pdf
- ❖ W. Fan J. Li, S. Ma, and N. Tang, and Y. Wu. *Graph pattern matching: From intractable to polynomial time*, VLDB, 2010.
- ❖ W. Fan, F. Geerts, and F. Neven. *Making Queries Tractable on Big Data with Preprocessing*, VLDB 2013