

CPT-S 415

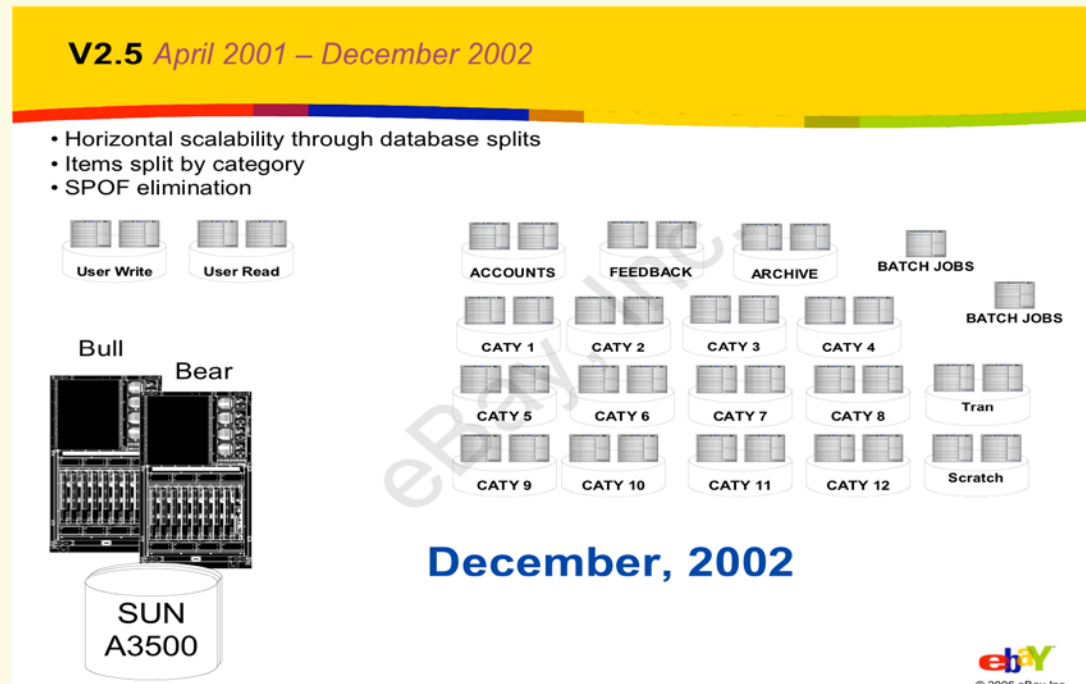
Big Data

Yinghui Wu

EME B45


Let's go back to Early-2000s...

- ✓ All the big players were heavyweight and expensive.
 - Oracle, DB2, Sybase, SQL Server, etc.
- ✓ Open-source databases were missing important features.
 - Postgres, mSQL, and MySQL.



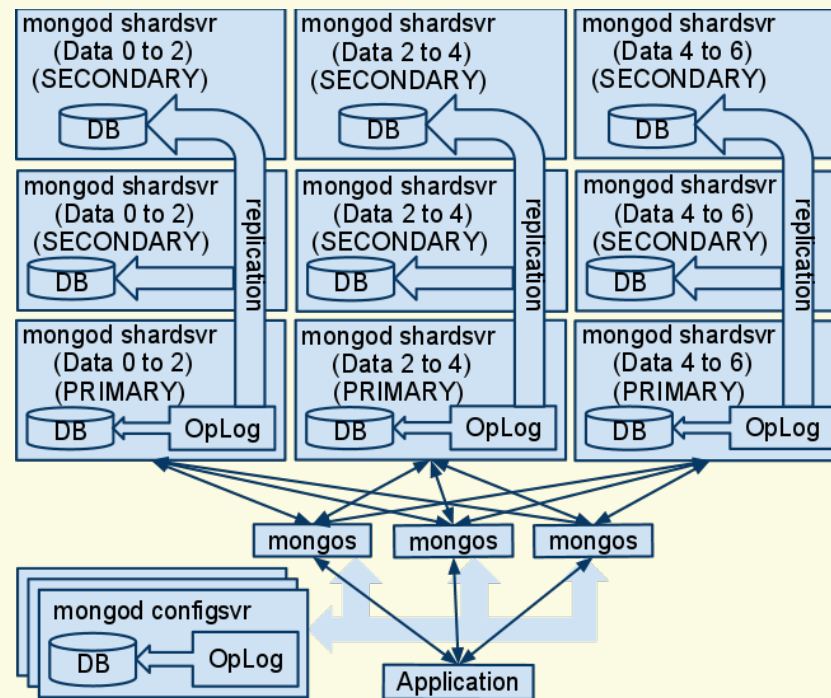
Mid-2000s

- ✓ MySQL + InnoDB is widely adopted by new web companies:
 - Supported transactions, replication, recovery.
 - Still must use custom middleware to scale out across multiple machines.
 - Memcache for caching queries.




Late-2000s

- ✓ NoSQL systems are able to scale horizontally:
 - Schemaless.
 - Using custom APIs instead of SQL.
 - Not ACID (i.e., eventual consistency)
 - Many are based on Google's BigTable or Amazon's Dynamo systems.

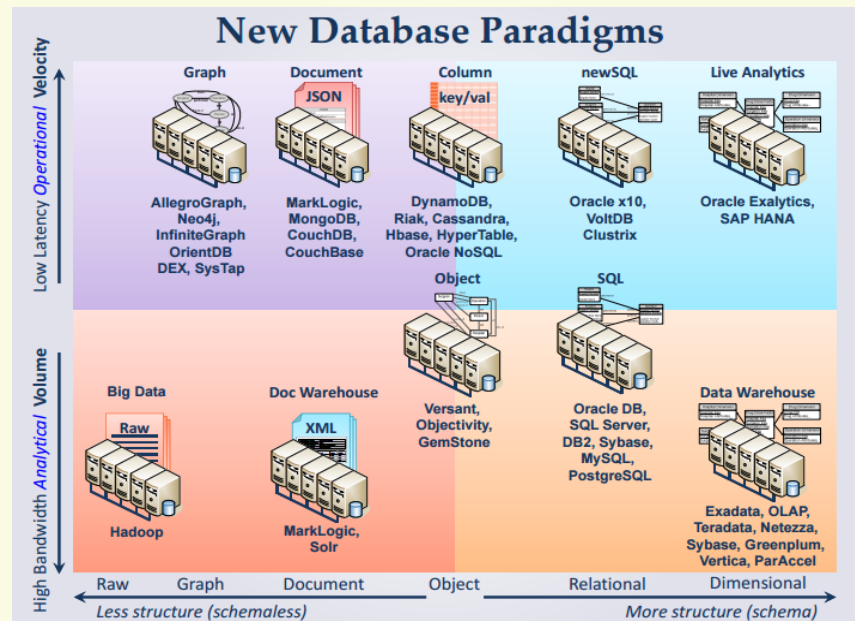


Early-2010s

- ✓ New DBMSs that can scale across multiple machines natively and provide ACID guarantees.
 - MySQL Middleware
 - Brand New Architectures
- ✓ “New SQL”



newSQL



old OLTP and old SQL

- ✓ An information system can be transactional (OLTP) or/and analytical (OLAP)

- ✓ **OLTP**

transa

- ve

- da

- ef

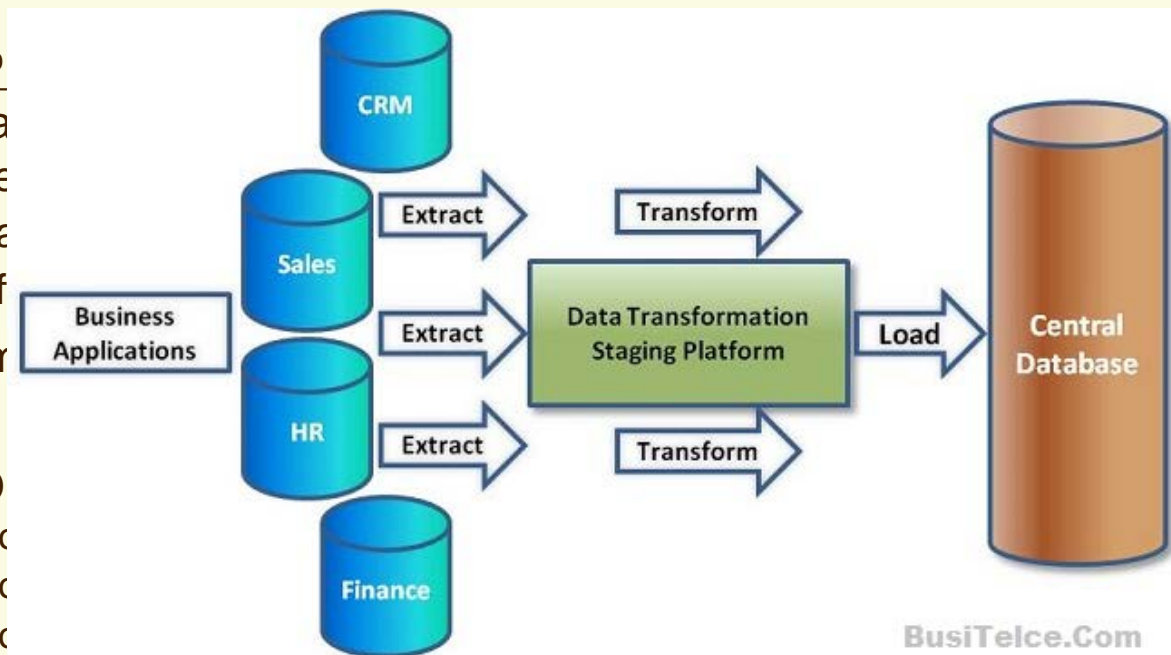
- ✓ schem

- ✓ Old O

- Co

- Co

- Co



port on-line

- ✓ Old SQL: techs, systems and vendors supporting old OLTP

New requirements (new OLTP)

- ✓ Web changes everything
- ✓ Large scale systems, with huge and growing data sets
 - ✓ 9M messages per hour in Facebook
 - ✓ 50M messages per day in Twitter
- ✓ Information is frequently generated by devices (cellphones, PDAs, sensors...)
-> **“Online”**
- ✓ High concurrency requirements, high-throughput ACID write ->
“Transaction”
- ✓ High Availability + Durability: core database requirements
- ✓ Need for high throughput
- ✓ Need for real-time analytics

Challenge

- ✓ Ingest the firehose in real time
- ✓ Process, validate, enrich and respond in Real-time
- ✓ Real-time analytics

- ✓ Options:
 - Old SQL - Legacy RDBMS vendors
 - NoSQL: give up SQL and ACID
 - NewSQL: SQL + ACID + new architecture

noSQL

- ✓ Give up SQL
- ✓ Give up ACID
 - Data accuracy
 - Funds transfer
 - Integrity constraints
 - Multi-record state
- ✓ noSQL fits
 - Non-transactional systems
 - Single record transactions that are commutative
- ✓ noSQL is not a good fit for
 - New OLTP: gaming, purchasing, order management, real-time analytical, etc

NewSQL: informal definition


- ✓ SQL is good
 - ✓ ACID is good
 - ✓ Figure out a way to make oldSQL perform
 - ✓ Make it scale like noSQL
 - ✓ Make it available
-
- ✓ “A DBMS that delivers the scalability and flexibility promised by NoSQL while retaining the support for SQL queries and/or ACID, or to improve performance for appropriate workloads.”

-- 451 Group

NewSQL: definition

- ✓ SQL as the primary mechanism for application interaction
- ✓ ACID support for transactions
- ✓ A non-locking concurrency control mechanism so real-time reads will not conflict with writes, and thereby cause them to stall.
- ✓ An architecture providing much higher per-node performance than available from the traditional "elephants"
- ✓ A scale-out, shared-nothing architecture, capable of running on a large number of nodes without bottlenecking
 - Michael Stonebraker

Traditional DBMS overheads



“Removing those overheads and running the database in main memory would yield orders of magnitude improvements in database performance”

NewSQL design principles

- ✓ SQL + ACID + performance and scalability through modern innovative software architecture
- ✓ Principle 1: minimizing or stay away from locking
- ✓ Principle 2: rely on main memory
- ✓ Principle 3: try to avoid latching
- ✓ Principle 4: cheaper solutions for HA

NewSQL design principles

- ✓ new solution other than low-level record level locking mechanism
 - Transaction processed in timestamp order with no locking (voltDB)
 - multisession concurrency control (nuoDB)
 - Optimistic concurrency control (Google)
 - Principle: minimizing or stay away from locking

NewSQL design principles


- ✓ new solution for buffer pool overhead
 - Main memory DBMS
 - Moderate case is tilted towards main memory
 - [Principle 2: rely on main memory](#)
- ✓ new solution to latching for shared data structures
 - new way to manage B-trees
 - Single-threading
 - [Principle 3: try to avoid latching](#)
- ✓ new solution for write-ahead logging
 - Built-in replication
 - [Principle 4: cheaper solutions for HA](#)



newSQL databases

NewSQL: categories

- ✓ New approaches: VoltDB, Clustrix, NuoDB
- ✓ New Storage engines: TokuDB, ScaleDB
- ✓ Transparent Clustering: ScaleBase, dbShards



voltDB

VoltDB

- ✓ VoltDB is an in-memory, horizontally scalable, ACID compliant, fast RDBMS
- ✓ Backed and architected by Michael Stonebraker
- ✓ An open source project
- ✓ Java + C/++
- ✓ Available in community and commercial editions

Technical Overview

- ✓ VoltDB tries to avoid the overhead of traditional databases
 - ✓ K-safety for fault tolerance
 - ✓ In memory operation for maximum throughput
 - reduce buffer management
 - ✓ Partitions operate autonomously and single-threaded
 - no latching or locking
- ✓ Built to horizontally scale

Technical Overview – Partitions (1/3)

- ✓ One **partition** per physical CPU core
 - Each physical server has multiple VoltDB partitions
- ✓ **Data** - Two types of tables
 - **Partitioned**
 - Single column serves as partitioning key
 - Rows are spread across all VoltDB partitions by partition column
 - Transactional data (high frequency of modification)
 - **Replicated**
 - All rows exist within all VoltDB partitions
 - Relatively static data (low frequency of modification)
- ✓ **Code** - Two types of **work** – both ACID
 - **Single-Partition**
 - All insert/update/delete operations within single partition
 - Majority of transactional workload
 - **Multi-Partition**
 - CRUD against partitioned tables across multiple partitions
 - Insert/update/delete on replicated tables

Technical Overview – Partitions (2/3)

✓ Single-partition vs. Multi-partition

select count(*) from orders where customer_id = 5
single-partition

select count(*) from orders where product_id = 3
multi-partition

insert into orders (customer_id, order_id, product_id) values (3,303,2)
single-partition

update products set product_name = 'spork' where product_id = 3
multi-partition

Partition 1		
1	101	2
1	101	3
4	401	2
1	knife	
2	spoon	
3	fork	

Partition 2		
2	201	1
5	501	3
5	502	2
1	knife	
2	spoon	
3	fork	

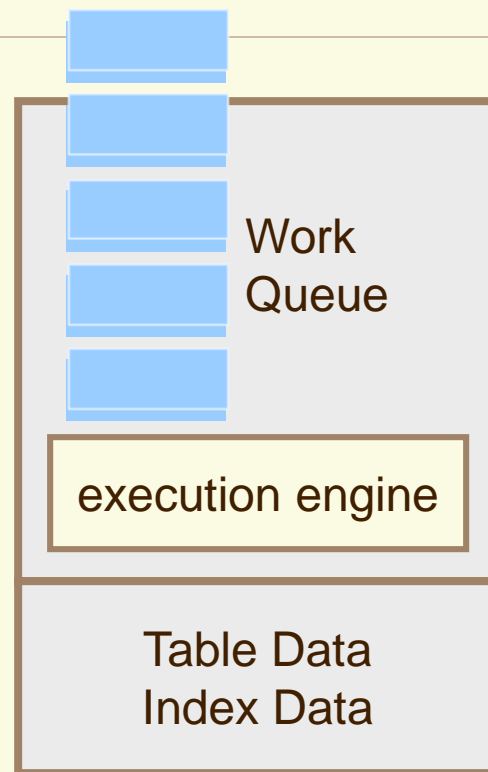
Partition 3		
3	201	1
6	601	1
6	601	2
1	knife	
2	spoon	
3	fork	

table orders : customer_id (partition key)
(partitioned) order_id
 product_id

table products : product_id
(replicated) product_name

Technical Overview – Partitions (3/3)

- ✓ Looking inside a VoltDB partition...
 - Each partition contains data and an execution engine.
 - The execution engine contains a queue for transaction requests.
 - Requests are executed sequentially (single threaded).



- Complete copy of all replicated tables
- Portion of rows (about 1/partitions) of all partitioned tables

Technical Overview – Compiling

- ✓ The database is constructed from
 - The schema (DDL)
 - The work load (Java stored procedures)
 - The Project (users, groups, partitioning)
- ✓ VoltCompiler creates application catalog
 - Copy to servers along with 1 .jar and 1 .so
 - Start servers

Schema


```
CREATE TABLE HELLOWORLD (  
  HELLO CHAR(15),  
  WORLD CHAR(15),  
  DIALECT CHAR(15),  
  PRIMARY KEY (DIALECT)  
);
```

Stored Procedures

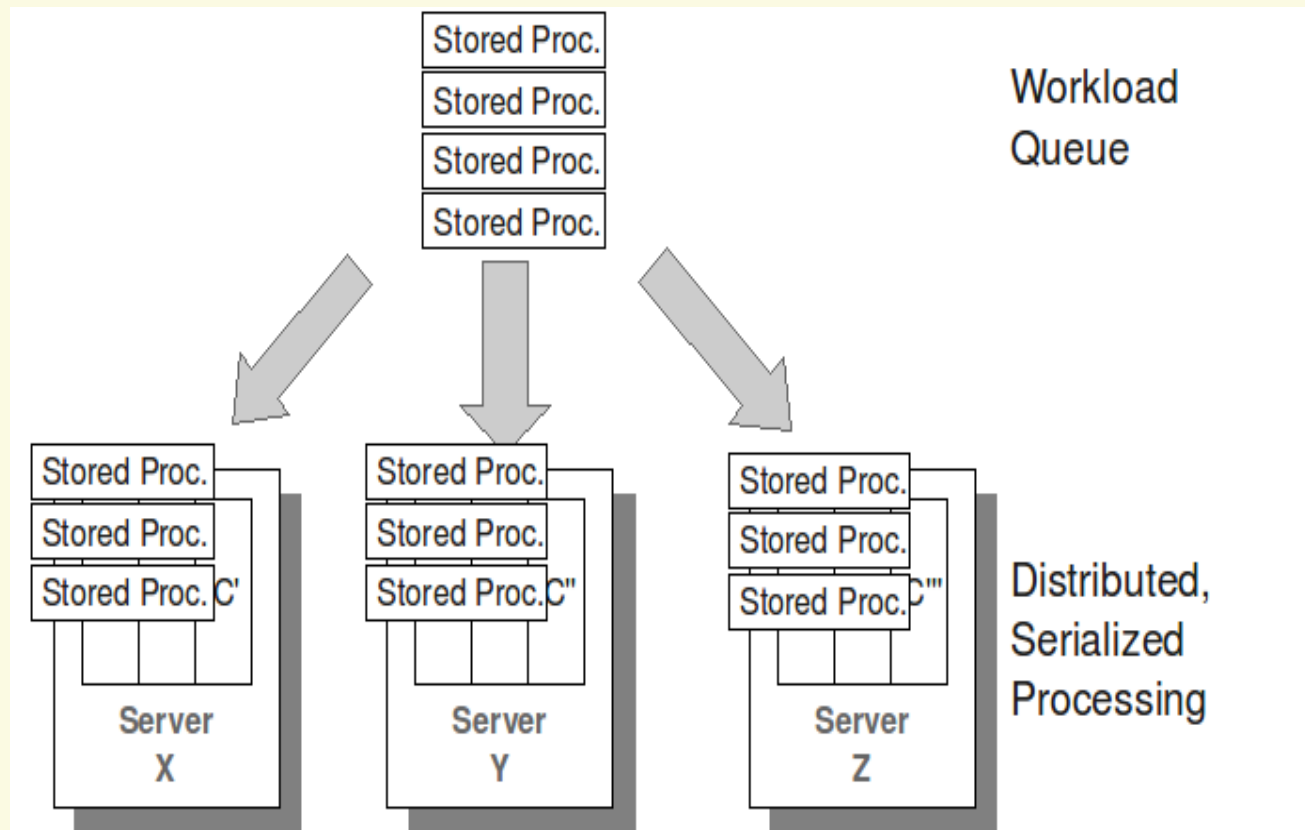
```
import org.voltadb. * ;  
  
@ProcInfo(  
  partitionInfo = "HE  
  singlePartition = t  
  
public final SQLStmt  
public VoltTable[] run
```

Project.xml

```
<?xml version="1.0"?>  
<project>  
  <database name='data'  
    <schema path='ddl.'  
      <partition table=''  
    </database>  
  </project>
```



Transaction Model



Procedures **outed** to, **ordered** and **run** at partitions

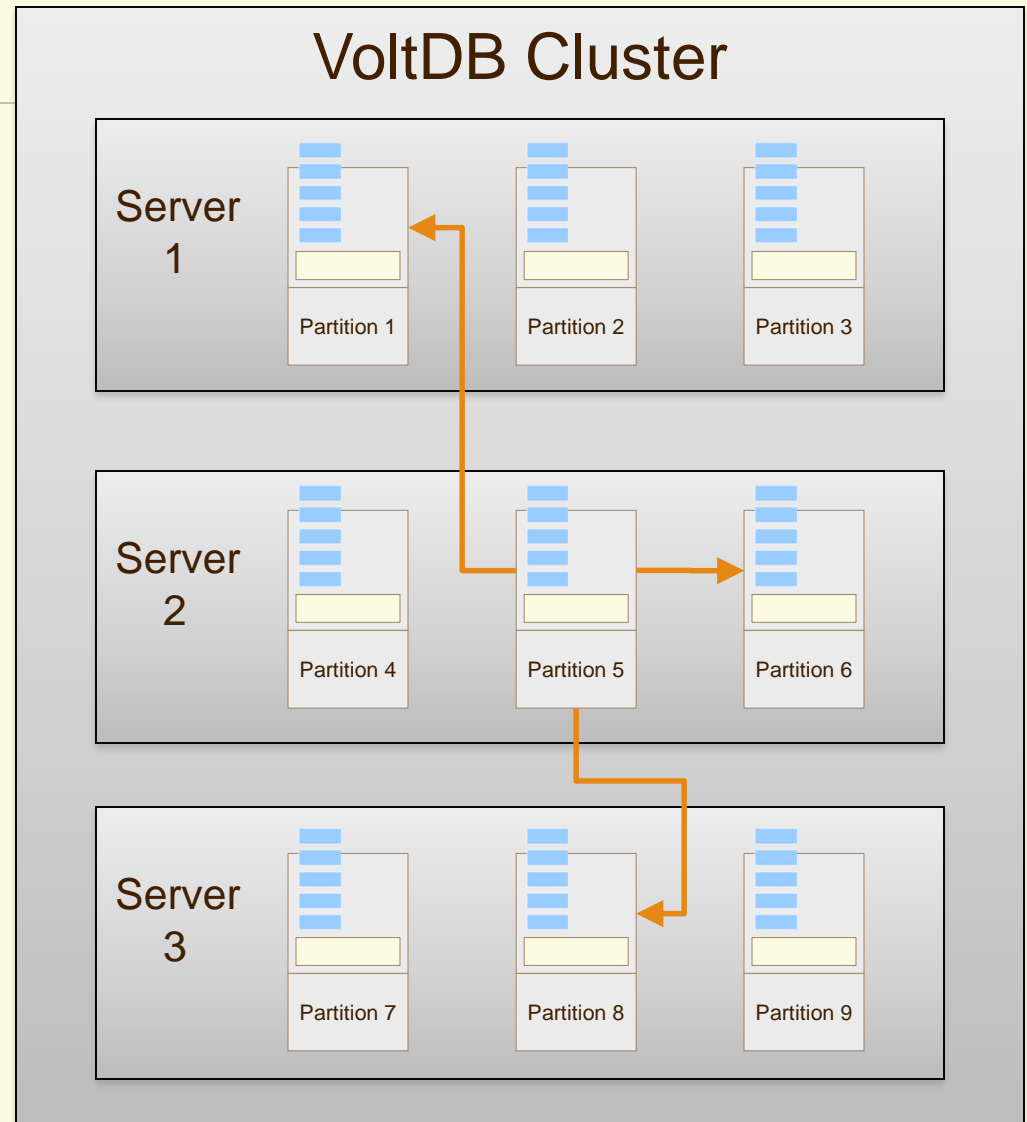
Transaction Execution

✓ Single partition transactions

- All data is in one partition
- Each partition operates autonomously

✓ Multi-partition transactions

- One partition distributes and coordinates work plans



Data Availability and Durability

✓ High Availability

- Data stored on server replicas (user configurable)
- Failover data redundancy
- No single point of failure

✓ Database Snapshots


- Simplifies backup/restore
- Scheduled, continuous, on demand
- Cluster-wide consistent copy of all data

✓ Command Logging

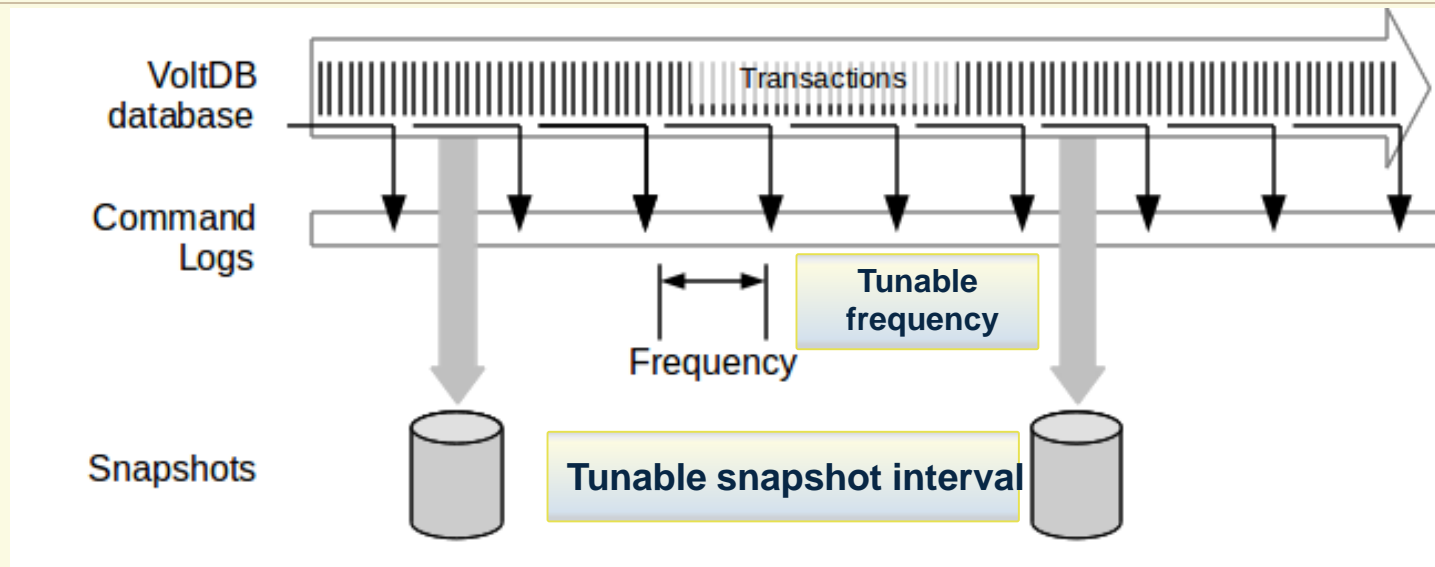
- Between Snapshots, every transaction is durable to disk

K-safety

- ✓ Duplicate database partitions for fault tolerance. K: # of replicas




Command Logging



- Synchronous logging provides highest durability at reduced performance
- Asynchronous logging best performance at reduced durability

Disaster Recovery

- ✓ Disk snapshots replicated via storage system
- ✓ Stream command logs from Primary to Replica
- ✓ Run from Replica on DR event, reverse on recovery



Lack of concurrency

- ✓ Single-threaded execution within partitions (single-partition) or across partitions (multi-partition)
- ✓ No concern about locking/dead-locks
 - great for “inventory” type applications
 - checking inventory levels
 - creating line items for customers
- ✓ Because of this, transactions execute in microseconds.
- ✓ However, single-threaded comes at a price
 - Other transactions wait for running transaction to complete
 - Useful for OLTP, not OLAP

Summary: NewSQL design principles

- ✓ SQL + ACID + performance and scalability through modern innovative software architecture
- ✓ Principle 1: minimizing or stay away from locking
- ✓ Principle 2: rely on main memory
- ✓ Principle 3: try to avoid latching
- ✓ Principle 4: cheaper solutions for HA

nuoDB

nuoDB

- ✓ nuoDB is an elastically scalable, ACID compliant, newSQL Database
- ✓ Backed and architected by Jim Starkley
- ✓ Runs on JVM
- ✓ Proprietary source project

NuoDB: Architecture



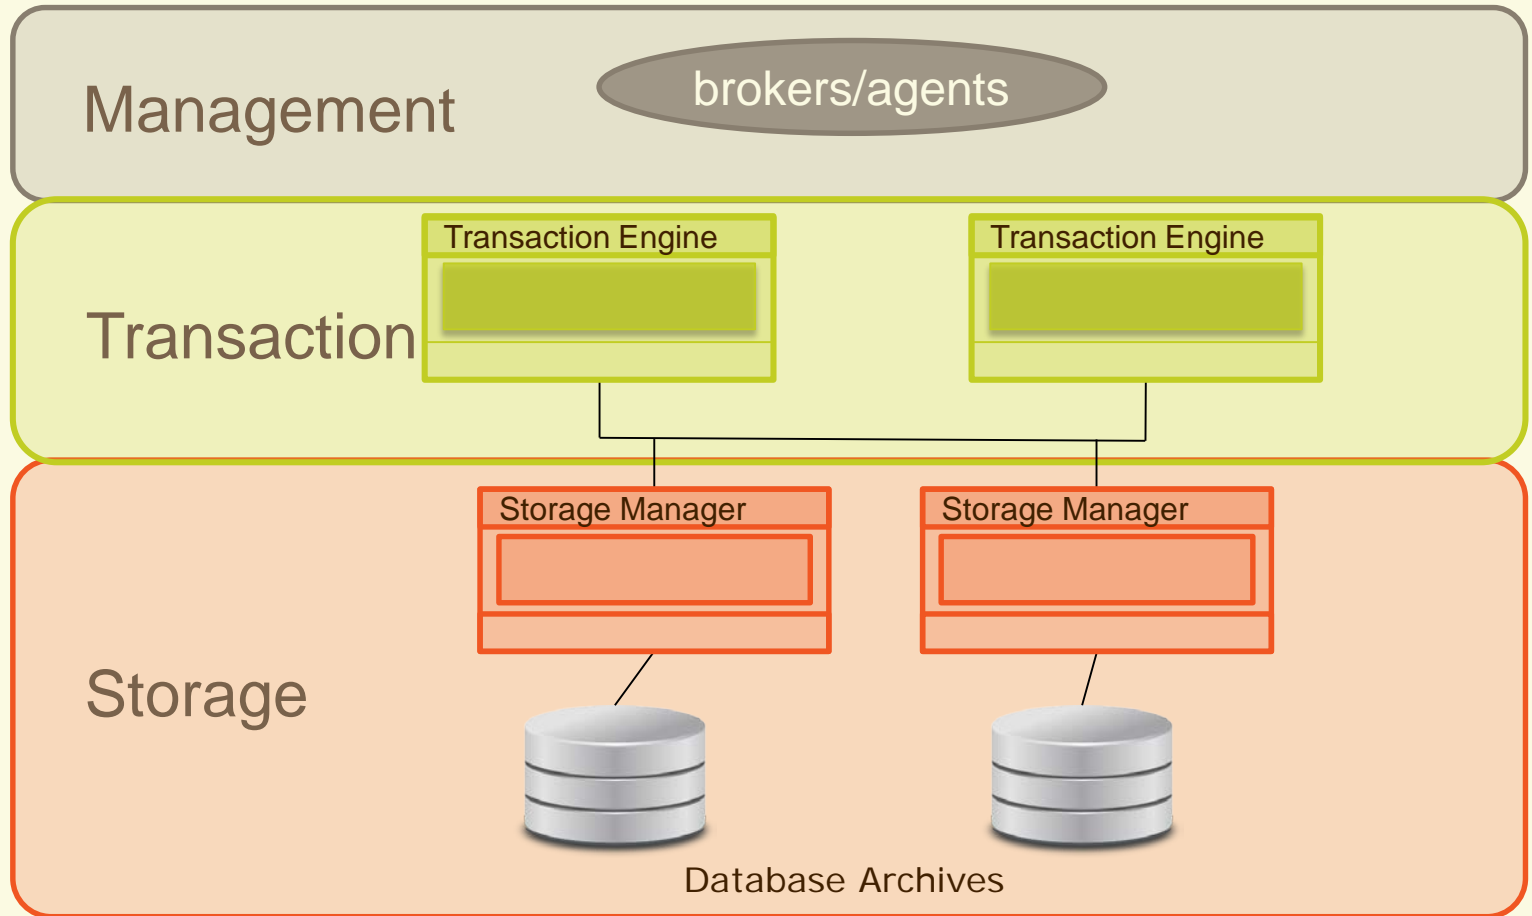
- ✓ Multi-tier Architecture
 - Transaction tier
 - Storage tier
 - Management tier
- ✓ Multi-Tenant
- ✓ Heavy use of memory
 - hot data stays in memory
 - Cold data in persistent store
- ✓ Object Oriented
 - Objects are atoms
- ✓ Asynchronous Messaging
- ✓ Partial, On-Demand replication
- ✓ MVCC - Concurrency

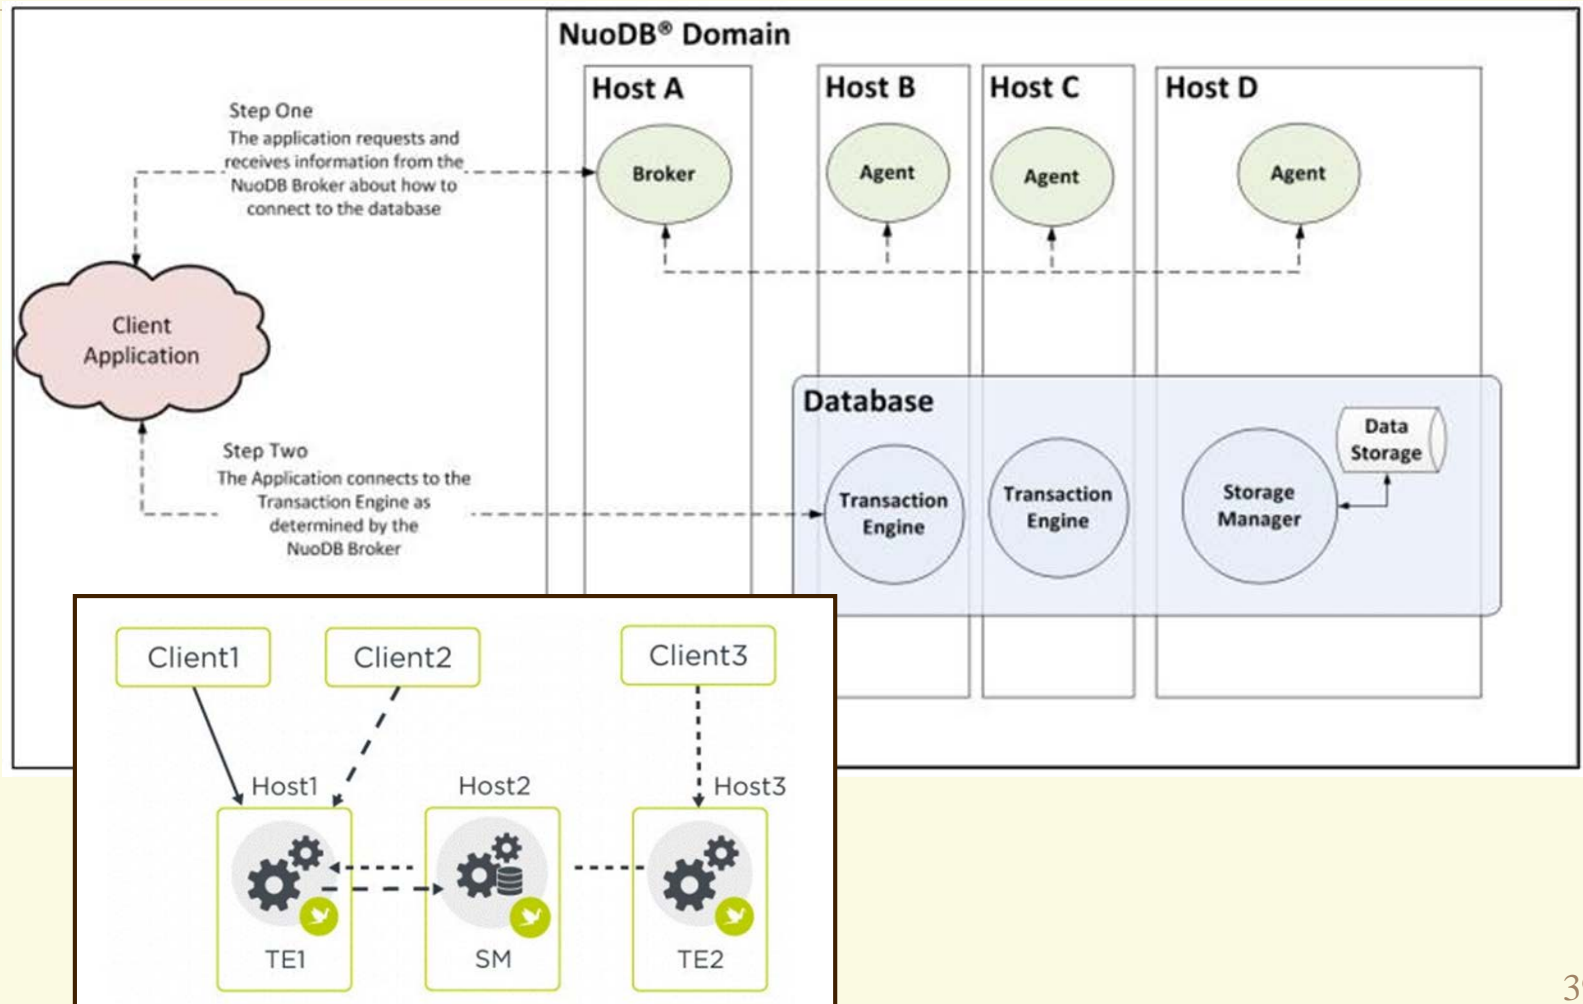
Technical Overview – Tiered Architecture

✓ Tiered Architecture

- Transactions: Transaction Engine
 - Parse, compile, optimize and execute SQL commands
 - Stores some information in memory locally
 - Map to locate the information
 - Any transaction engine can service any piece of information regardless of where it resides
 - Adding transaction engines -> More throughput
- Storage: Storage manager
 - Can run on HDFS or Amazon S3
 - Adding storage manager -> more resistance to failure
- Management: Agents and Brokers
 - (Re)starting transaction engines and storage managers
 - Collect statistics from them
 - Clients connect to TE via Brokers

Multi-tier architecture





Conclusions

- ✓ NewSQL is an established trend with a number of options
- ✓ Hard to pick one because they're not on a common scale
No silver bullet
- ✓ Growing data volume requires ever more efficient ways to store and process it

oldSQL, NoSQL and NewSQL: pros/cons

✓ OldSQL

- +: proven track record and standard SQL, ACID
- +: rich client ecosystem
- +: established
- -: not a scalable
- -: complex

✓ NoSQL

- +: higher availability
- +: support
- -: fundamental
- -: require a

✓ NewSQL

- +: stronger
- +: familiar
- +: richer architecture
- +: leverage
- -: no NewSQL is as general-purpose as traditional SQL systems
- -: in-memory architecture may hinder the application for large dataset
- -: partial access to the rich tooling of traditional SQL

